

# CAD Project I: Classical Machine Learning

---

KAOUTHER MOUHEB

RACHIKA ELHASSNA HAMADACHE

MAIA 2021-2023



# Outline

---

1. Proposal Analysis


2. Implementation and Design

1. Preprocessing
2. Feature Extraction
3. Sampling
4. Parameter Fine-tuning
5. Feature Selection
6. Model Selection

3. Experimental Results and Evaluation

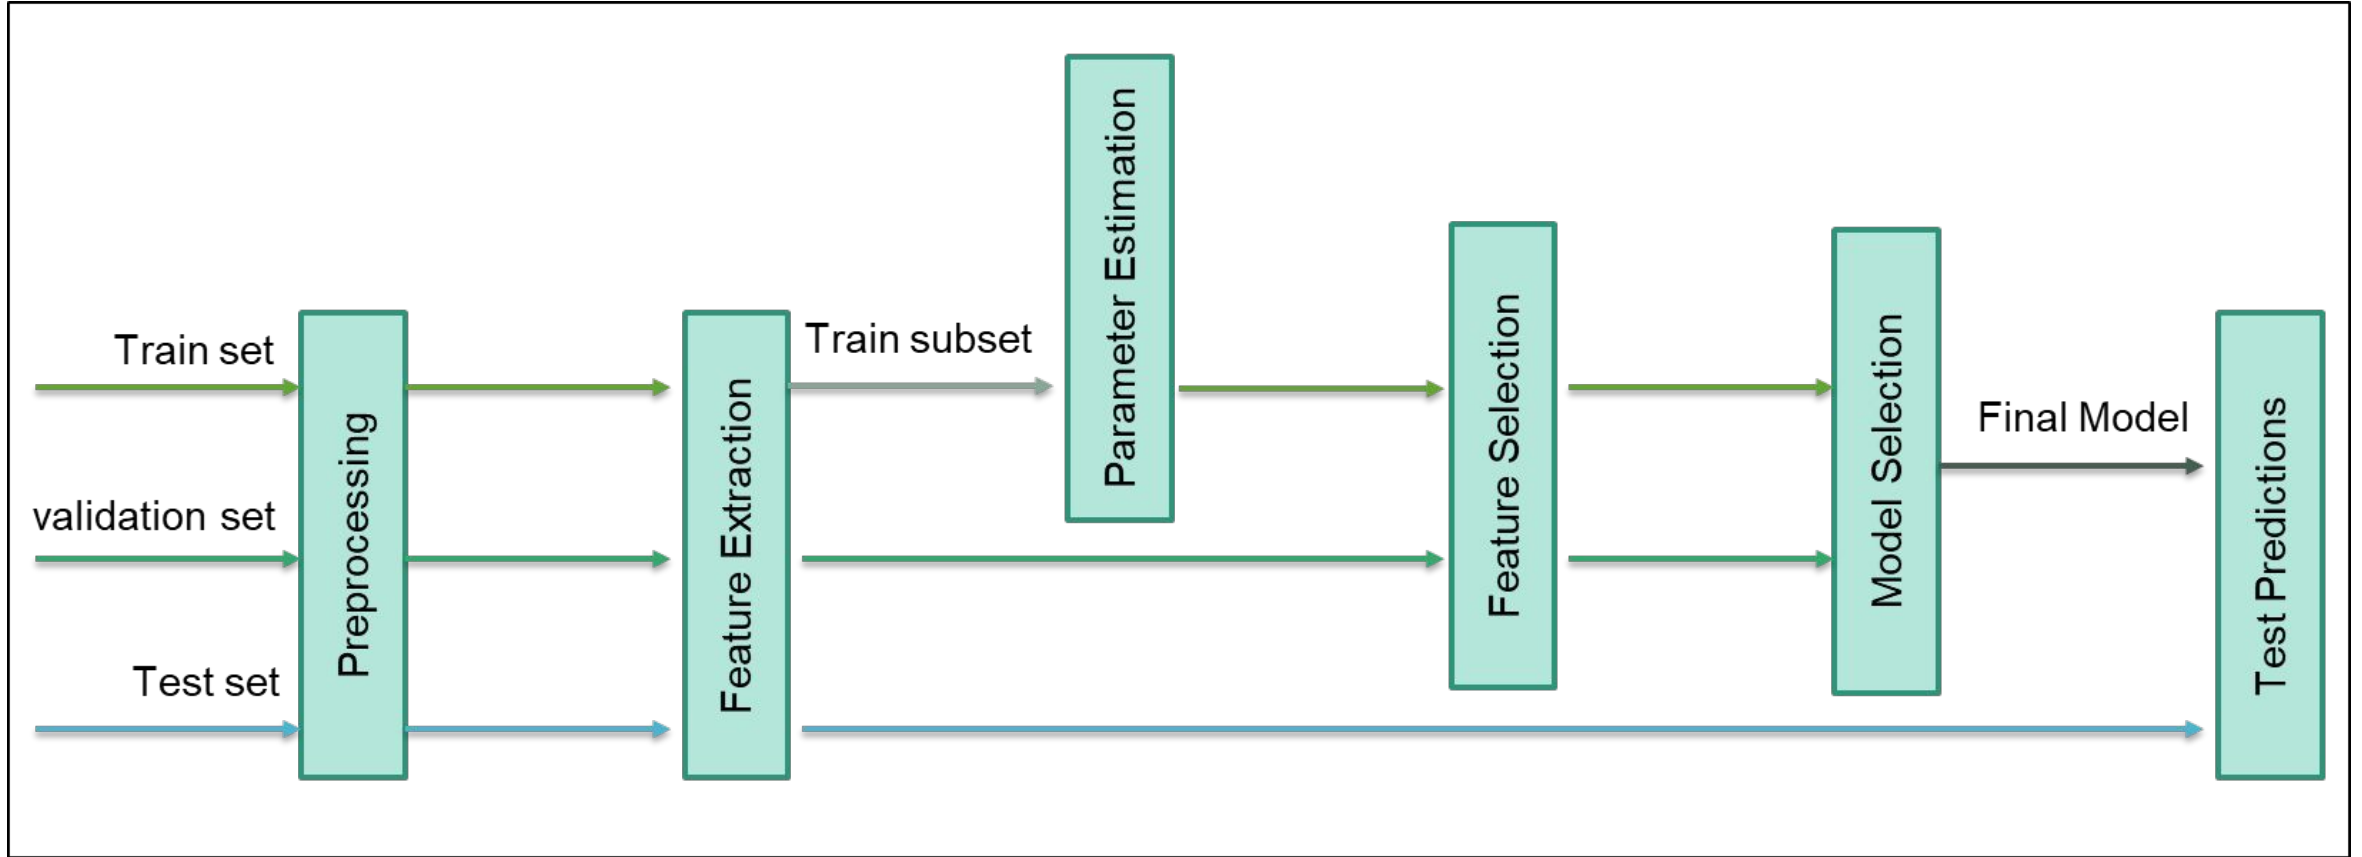
4. Bag Of Visual Words

5. Conclusion




# Proposal Analysis

# Proposed Pipeline



Proposed Pipeline



Preprocessing

# Hair removal

---





Figure 2: Hair removal pipeline

# Vignette removal

---



Vignette removal pipeline


# Color Constancy

A *color compensation* technique to reduce the influence of the acquisition setup on the color features extracted from the images. It uses the Minkowski norm to estimate the color of the illuminant.

The illuminant  $e$  is obtained as follows:

$$\left( \frac{\int (I_c(\mathbf{x}))^p d\mathbf{x}}{\int d\mathbf{x}} \right)^{1/p} = k e_c$$

Where  $I_c$  denotes the  $c^{\text{th}}$  component of the image,  $k$  a normalization factor,  $p$  the degree of the norm.

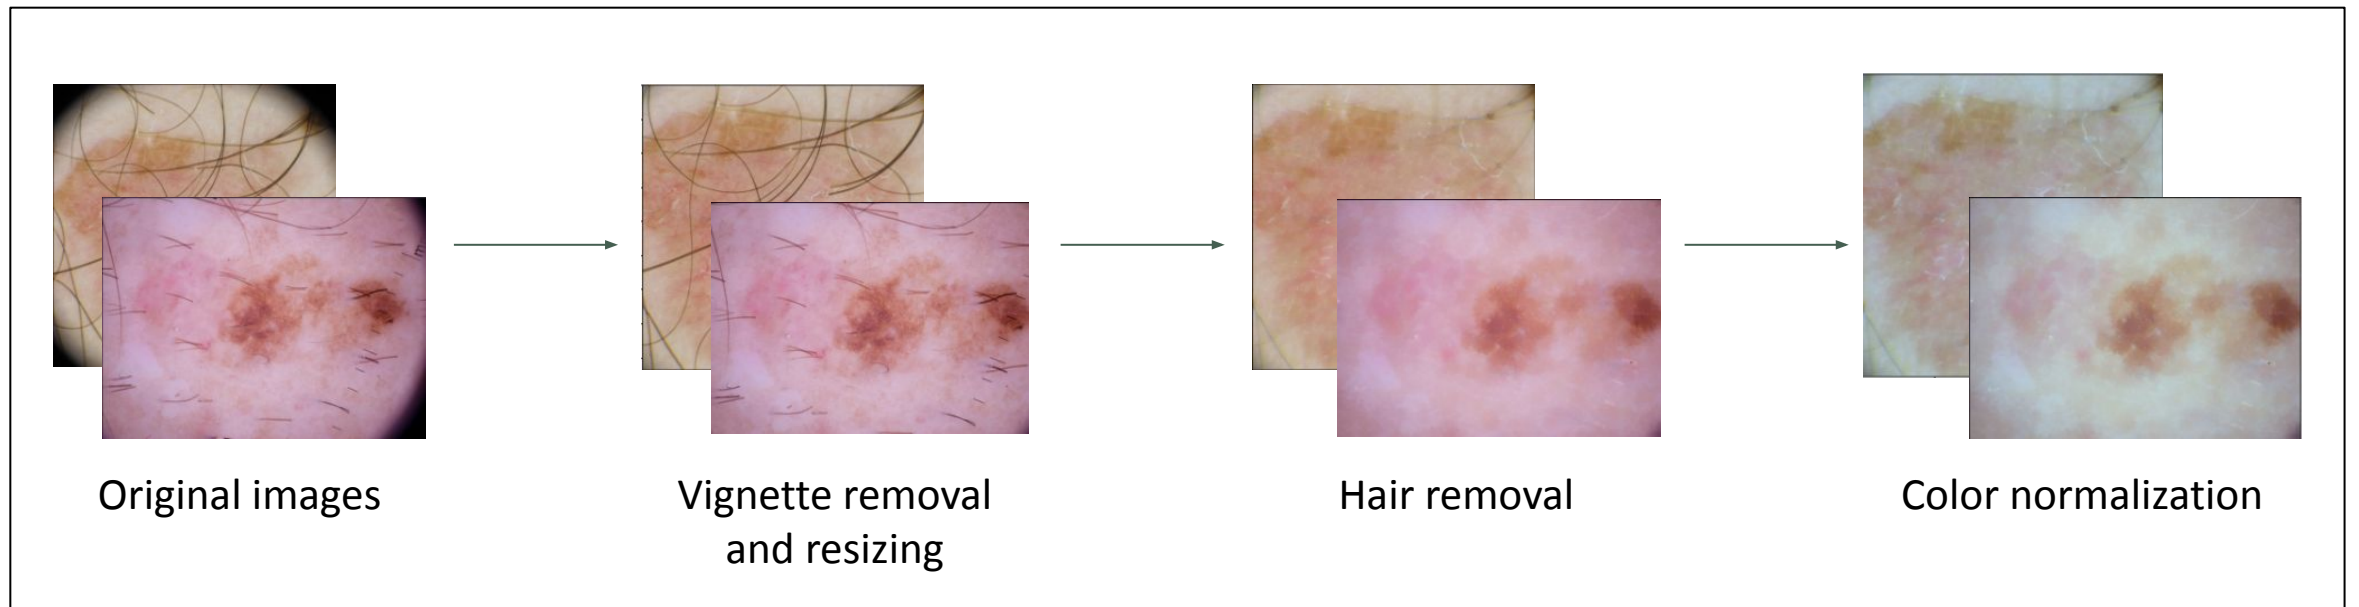


Color Normalization Example




# Full Preprocessing Pipeline

---



Preprocessing pipeline



# Design and Implementation

# Feature Extraction

# Color features

1. Color variegation : quantified by the normalized standard deviation of the red, green and blue components of the image.
2. Color moments : 4 measures characterizing the color distribution of the images : mean, standard deviation, skewness and variance.
3. Color histograms : histogram of each channel of the 3 color spaces : RGB, HSV and Lab

# Texture features

1. LBP : Local binary patterns with  $P=16$  and  $R=2$ .
2. Haralick : texture features based on the GLCM matrix and computing 13 measures.
3. GLCM : 5 main properties computed for 4 different angles with unit step : correlation, homogeneity, contrast, energy and dissimilarity

# Parameter Fine-tuning

# Parameter Estimation

---

- For challenge 1, a subset of 4000 images (2000 per class) was sampled to perform fine-tuning.
- All features are used for this step.
- GridSearchCV was used with 5 folds to estimate the Following parameters:

## **Random Forest and Extra Trees:**

- Maximum depth
- Number of Estimators
- Criterion

## **XGBoost:**

- Maximum depth
- Number of Estimators
- Learning rate


## **SVM:**

- Penalization parameter C.

## **K-NN:**

- The number of neighbors k.



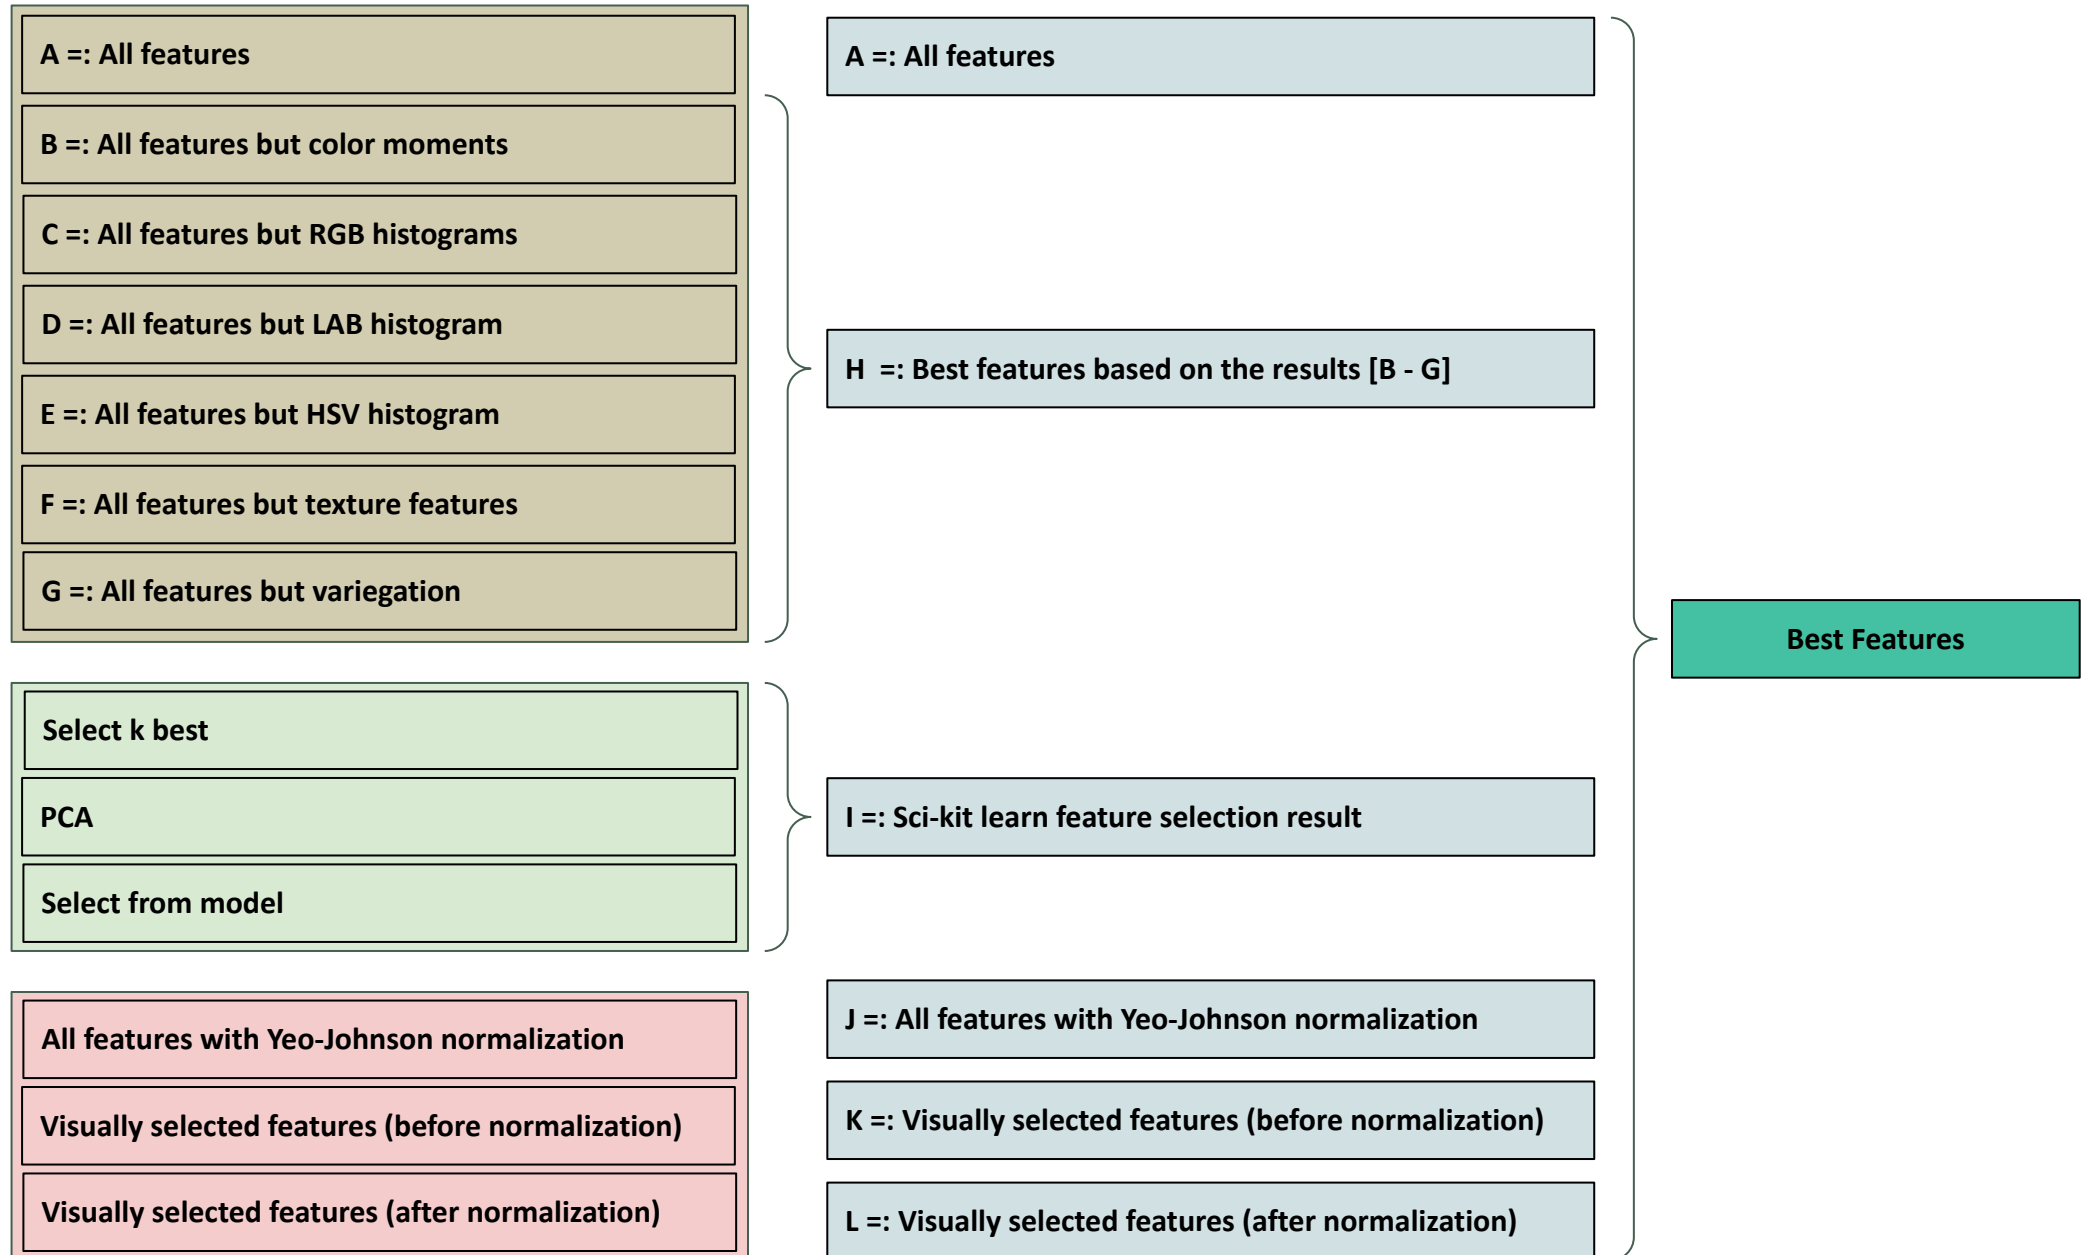


Challenge 1: Nevus Versus Others

# Feature Selection



**Feature Selection Pipeline (For each classifier)**



# Model Selection

# Challenge 1

---

The best model is selected based on the following evaluation metrics:

$$\kappa = \frac{2 \times (TP \times TN - FN \times FP)}{(TP + FP) \times (FP + TN) + (TP + FN) \times (FN + TN)}$$

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$


		True Labels	
		Nevus	Other
Predicted Labels	Nevus	TP	FP
	Other	FN	TN

# Challenge 1 (Cont'd)

---

The following models were compared:

- Nearest Neighbor (Baseline)
- K-Nearest Neighbors
- Support Vector Machines
- Gradient Boosting (XGBoost)
- Extra Trees
- Random Forest
- Stacking ensemble (SVM + XGboost + Extra trees with Logistic Regression as meta classifier)
- Averaging ensemble (SVM + XGboost + Extra trees)
- Majority voting ensemble (SVM + XGboost + Extra trees)
- Ensembling the pre-trained models with averaging
- Ensembling the pre-trained models with majority voting



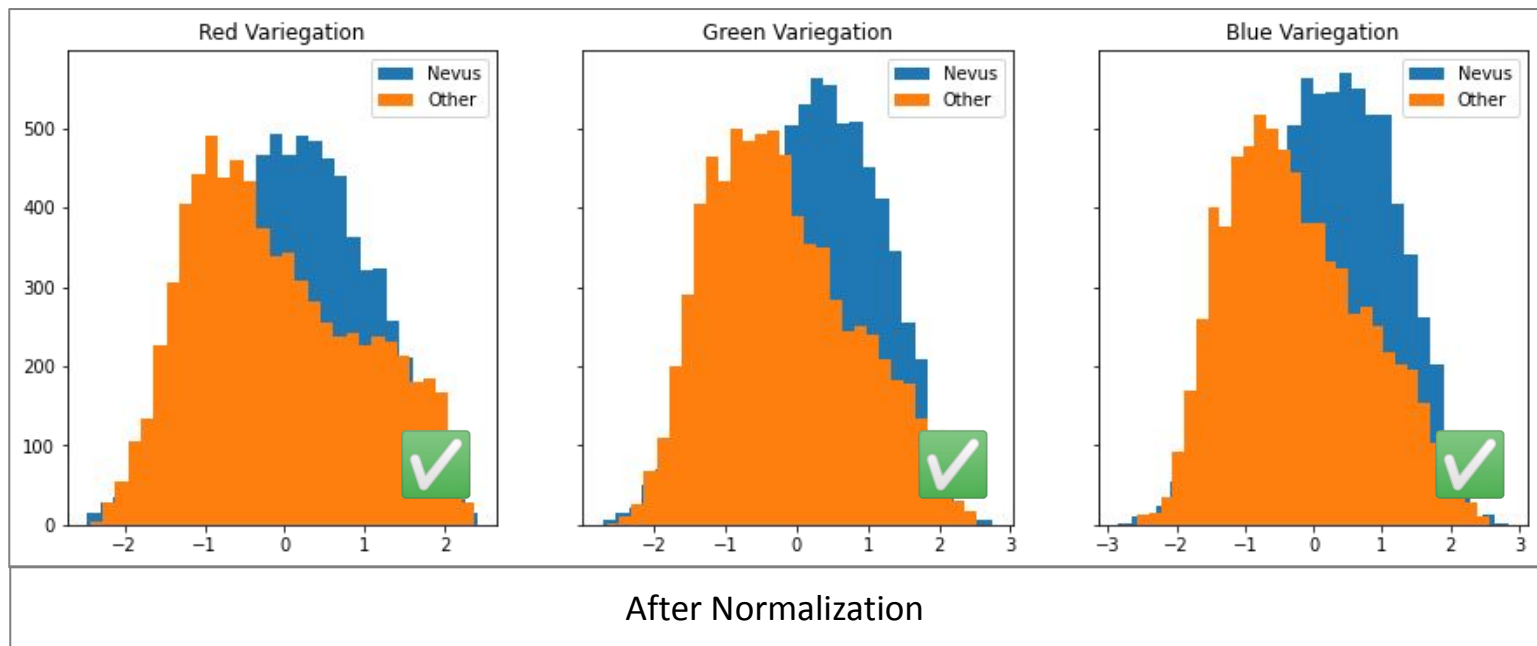
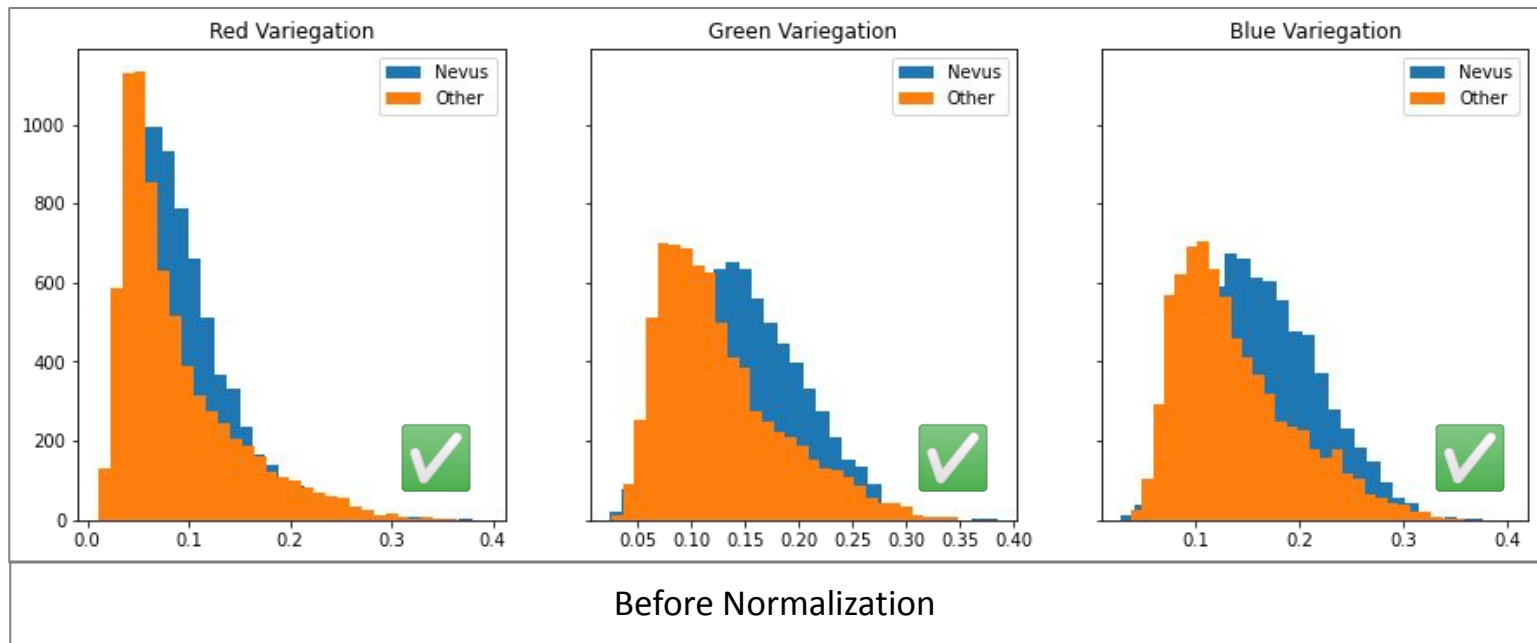
Experimental Results

# Feature Visualization

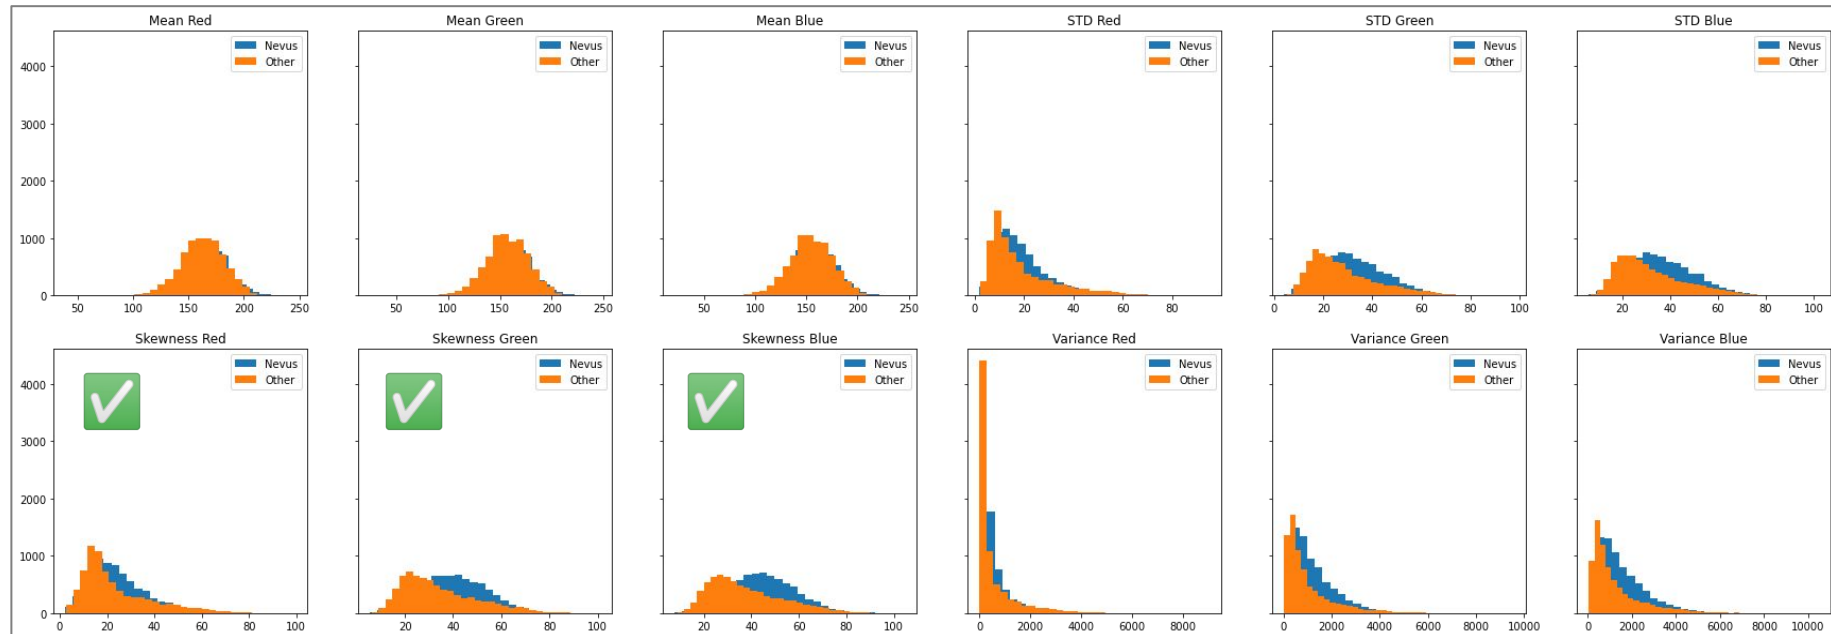
---

With Matplotlib

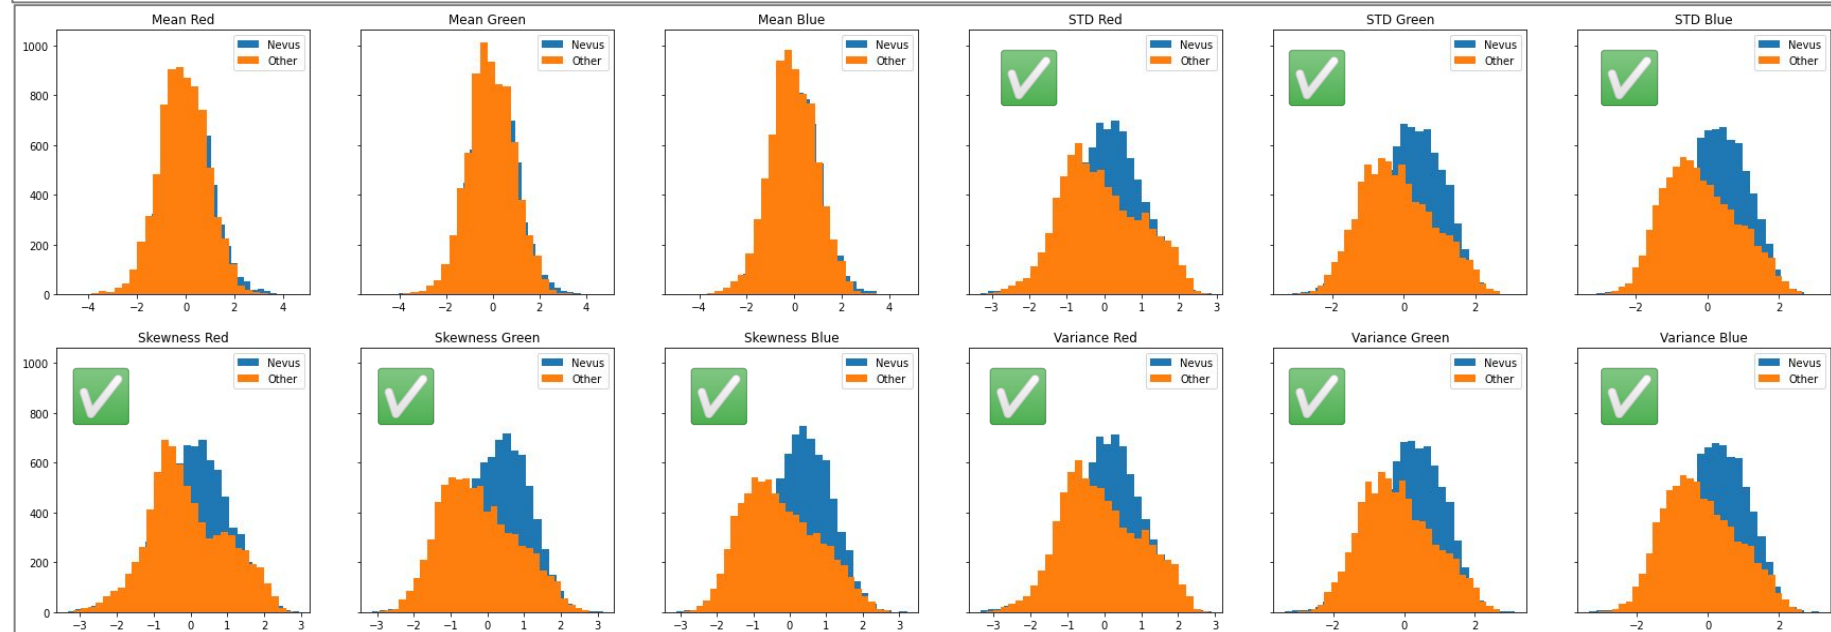
# Variegation



# Color Moments



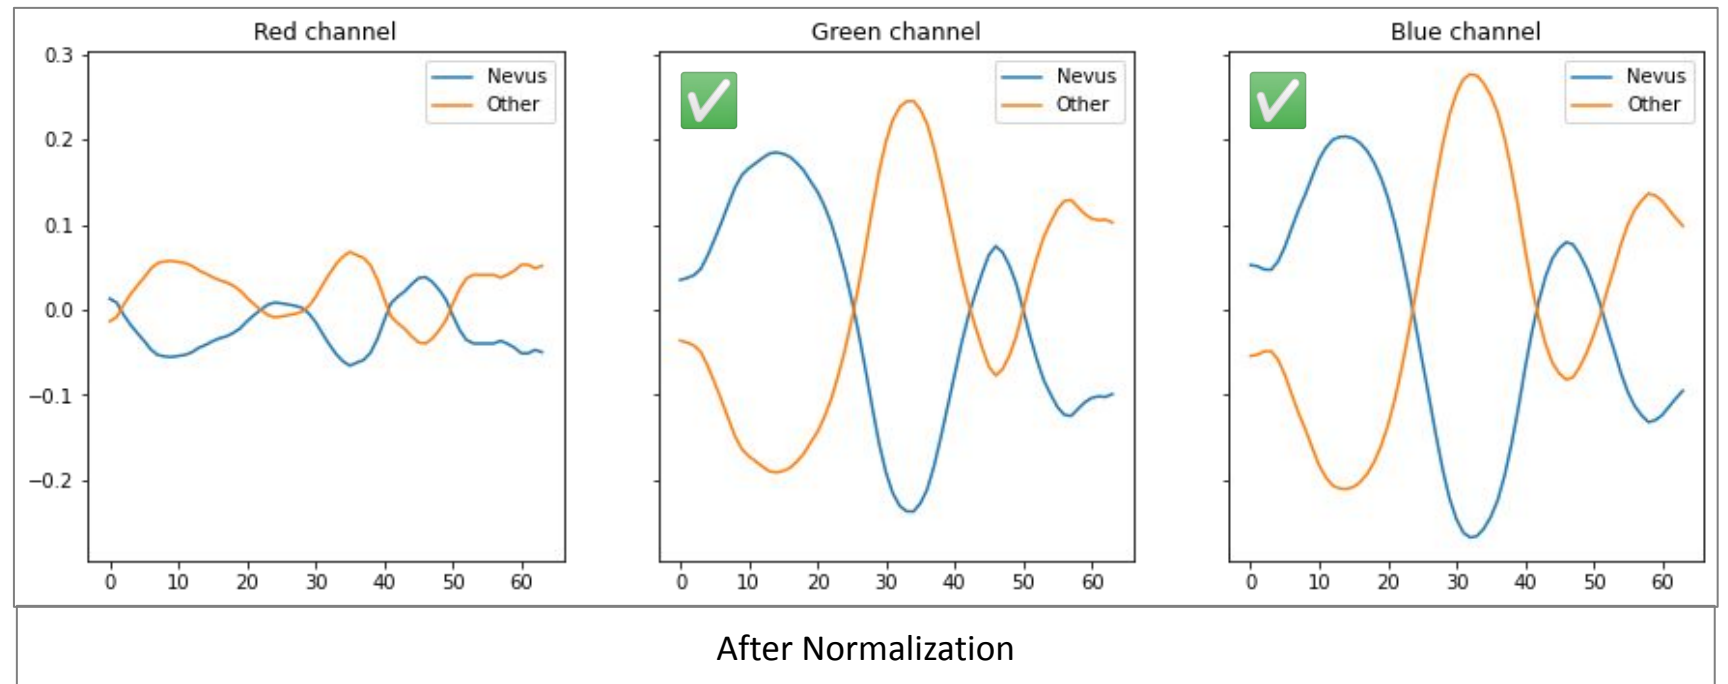
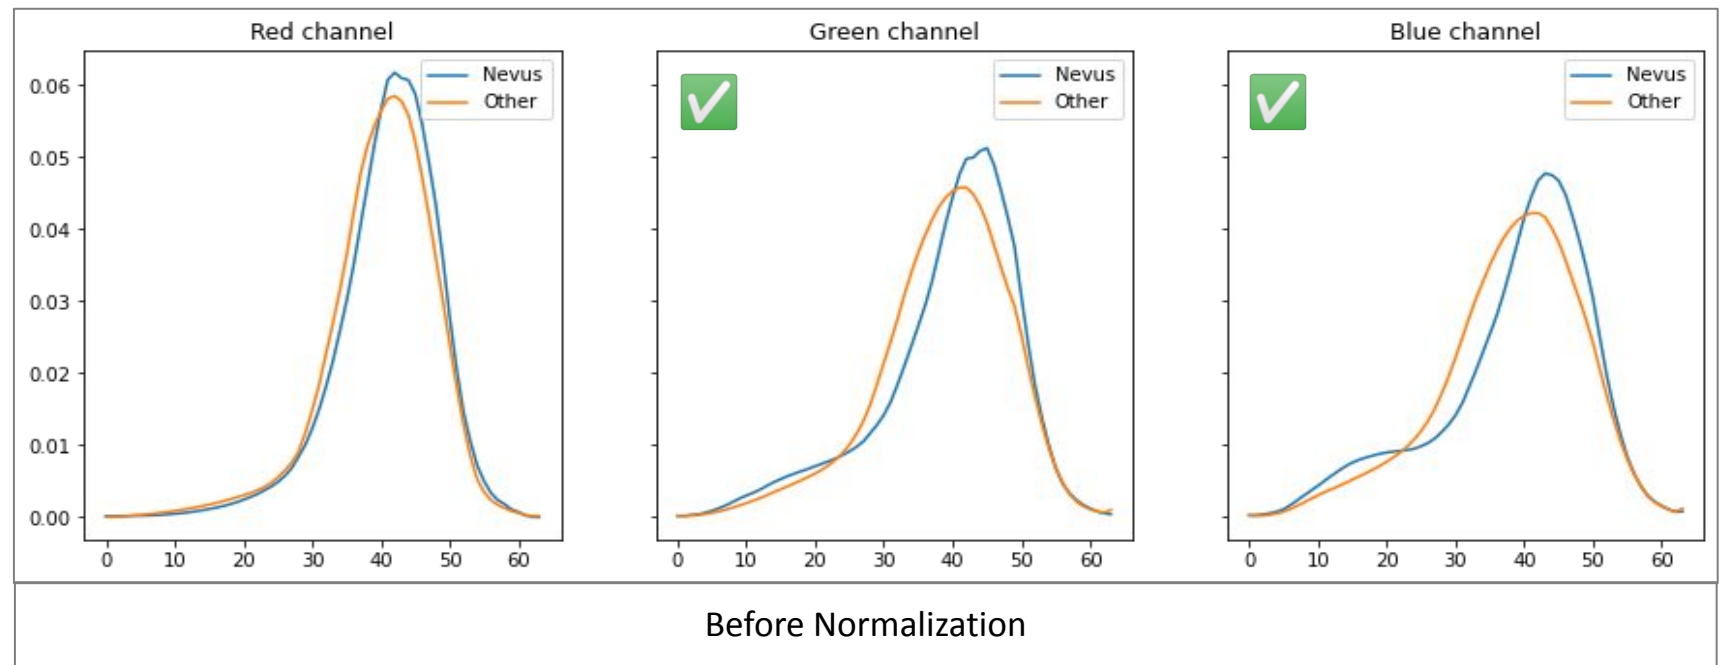
Before Normalization



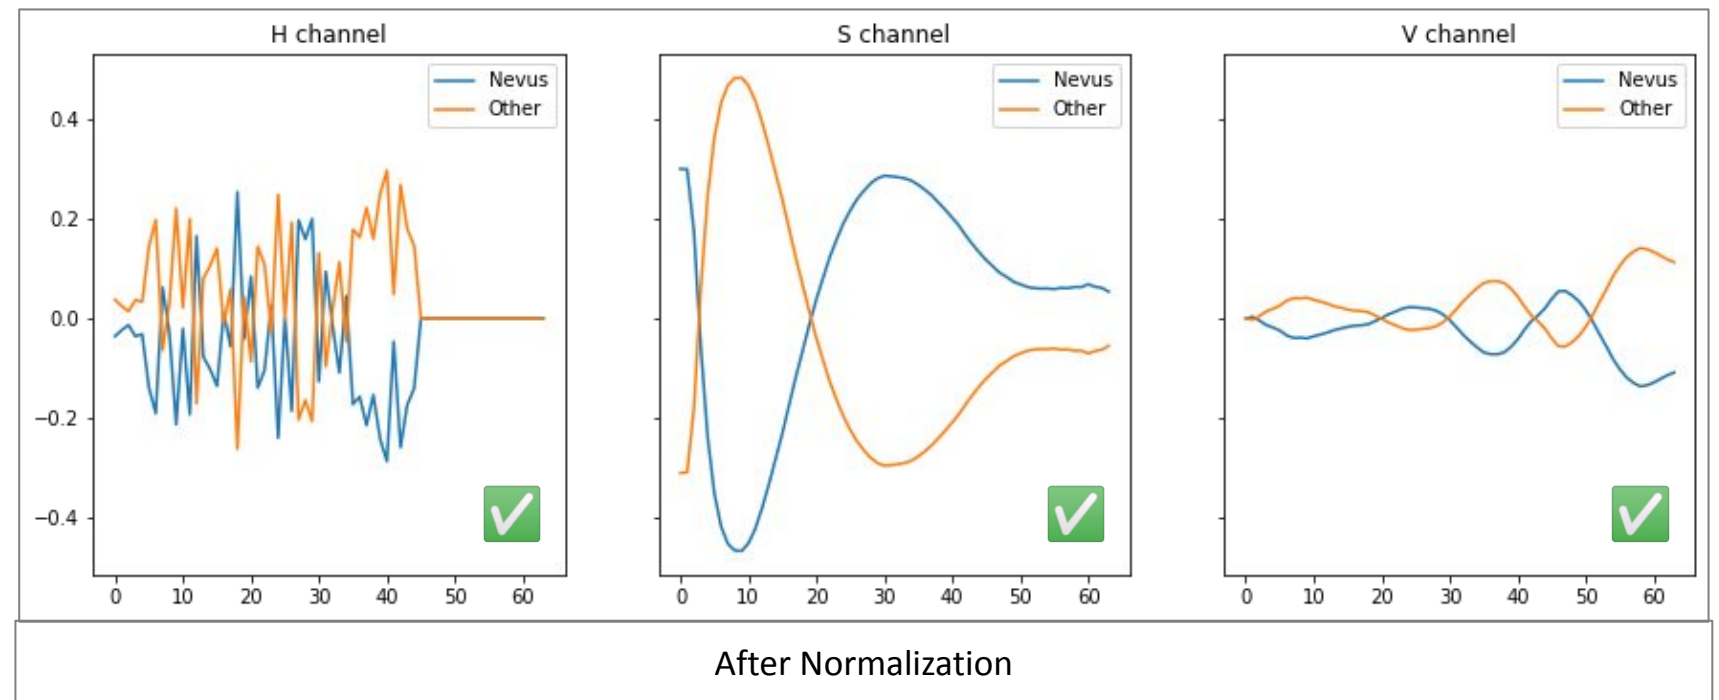
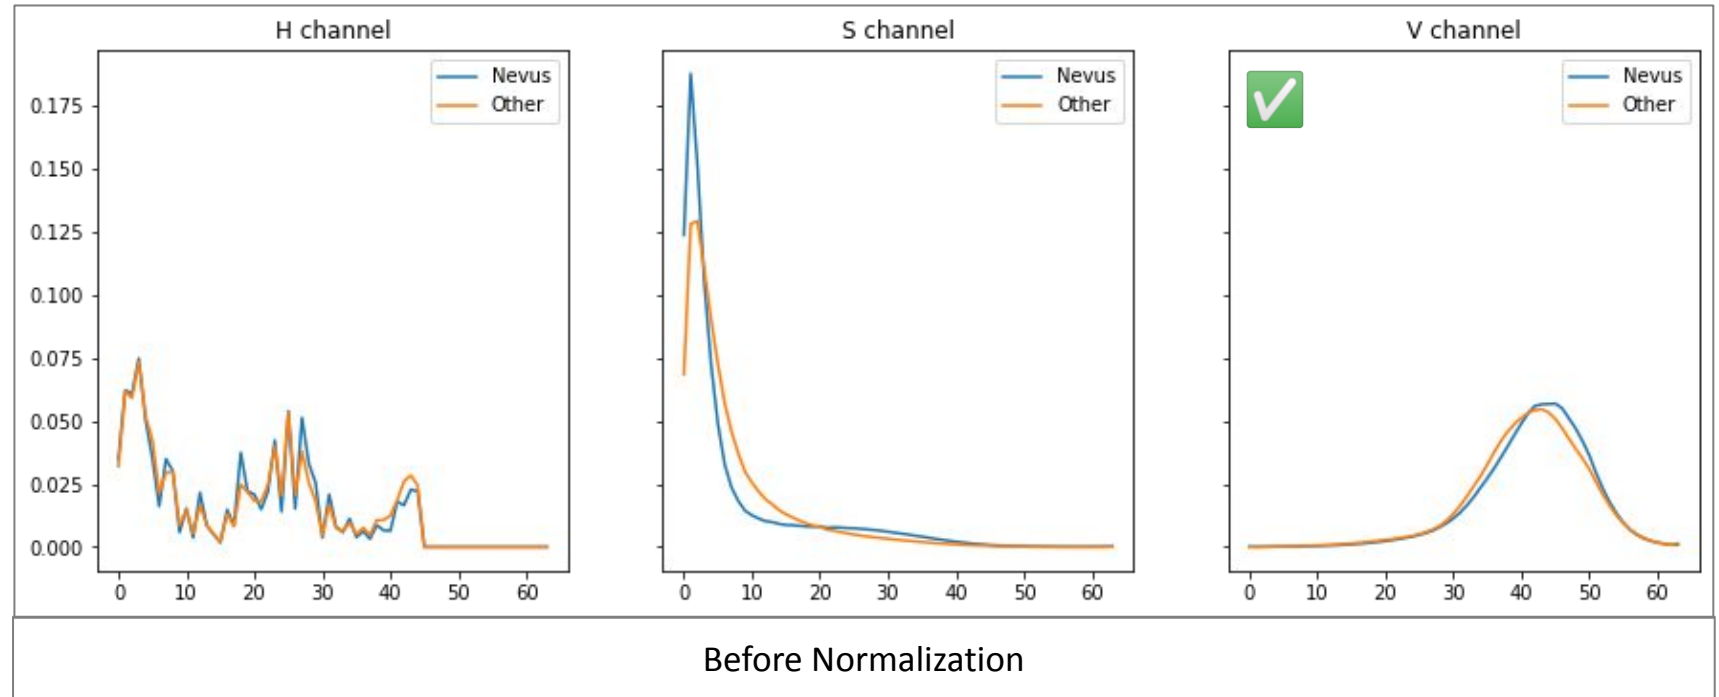
After Normalization





# Mean RGB Histograms For Each Class



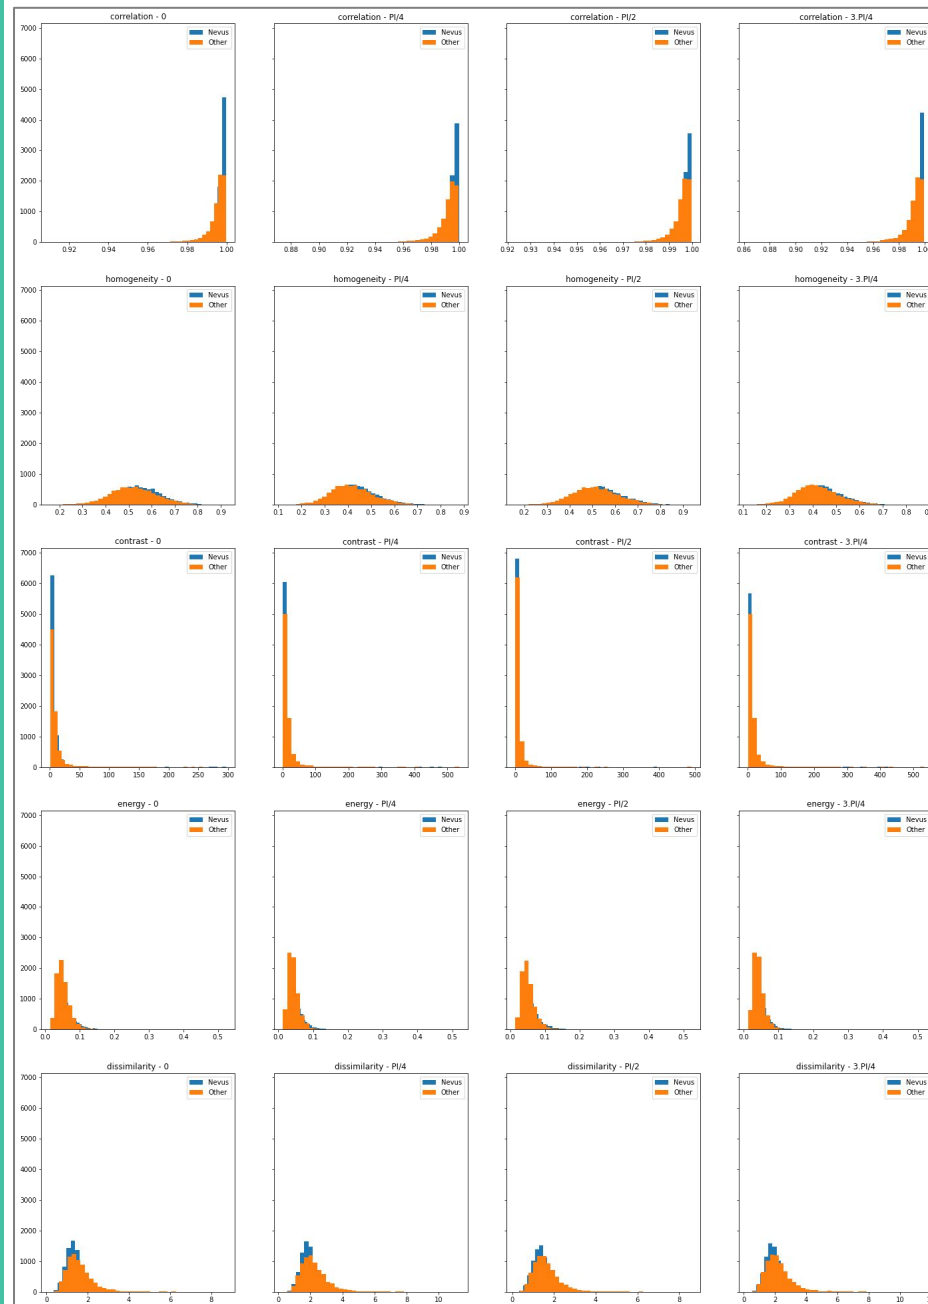
# Mean HSV Histograms For Each Class



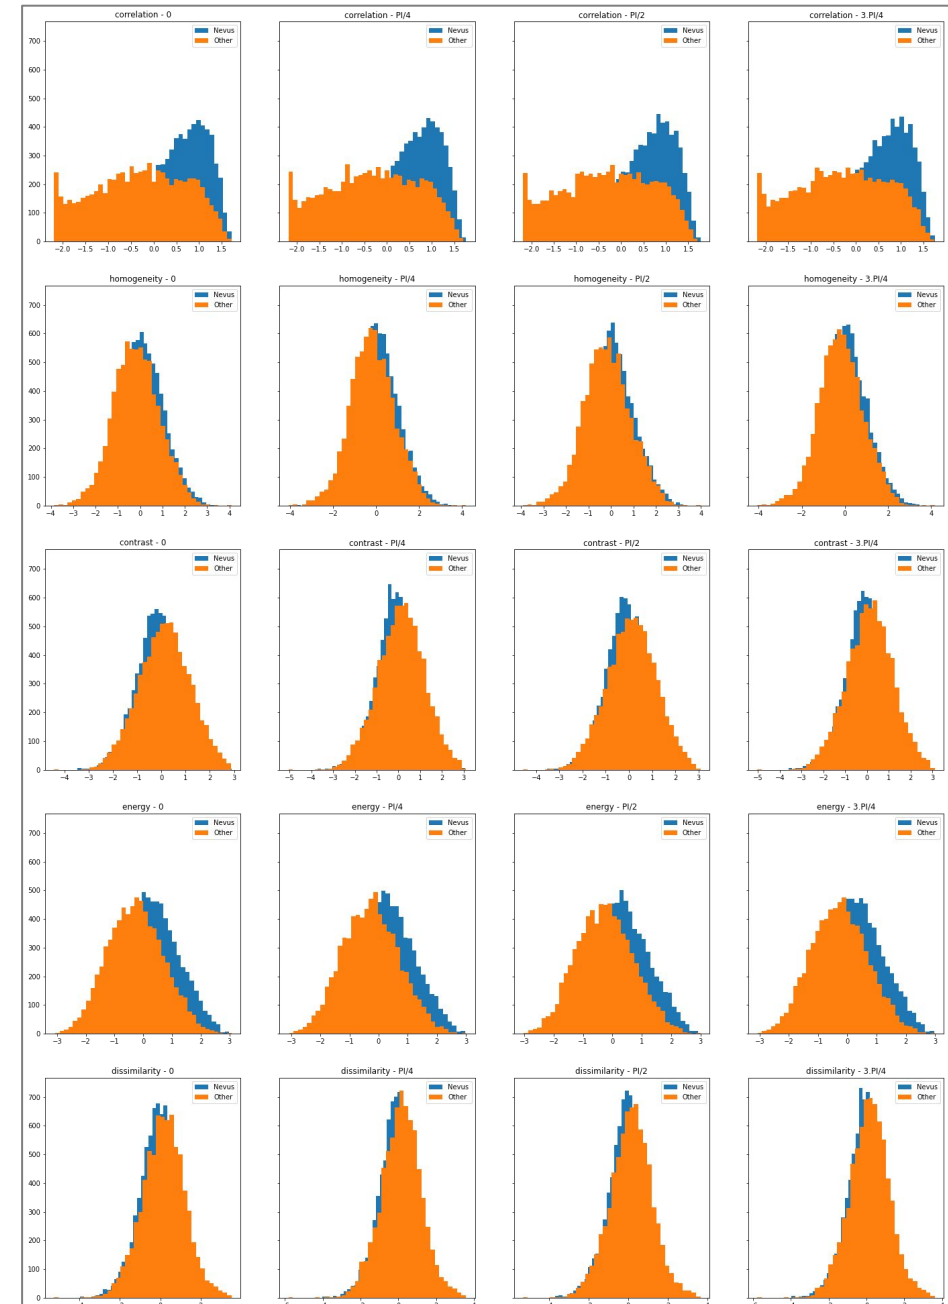
# Mean LAB Histograms For Each Class



# GLCM Features

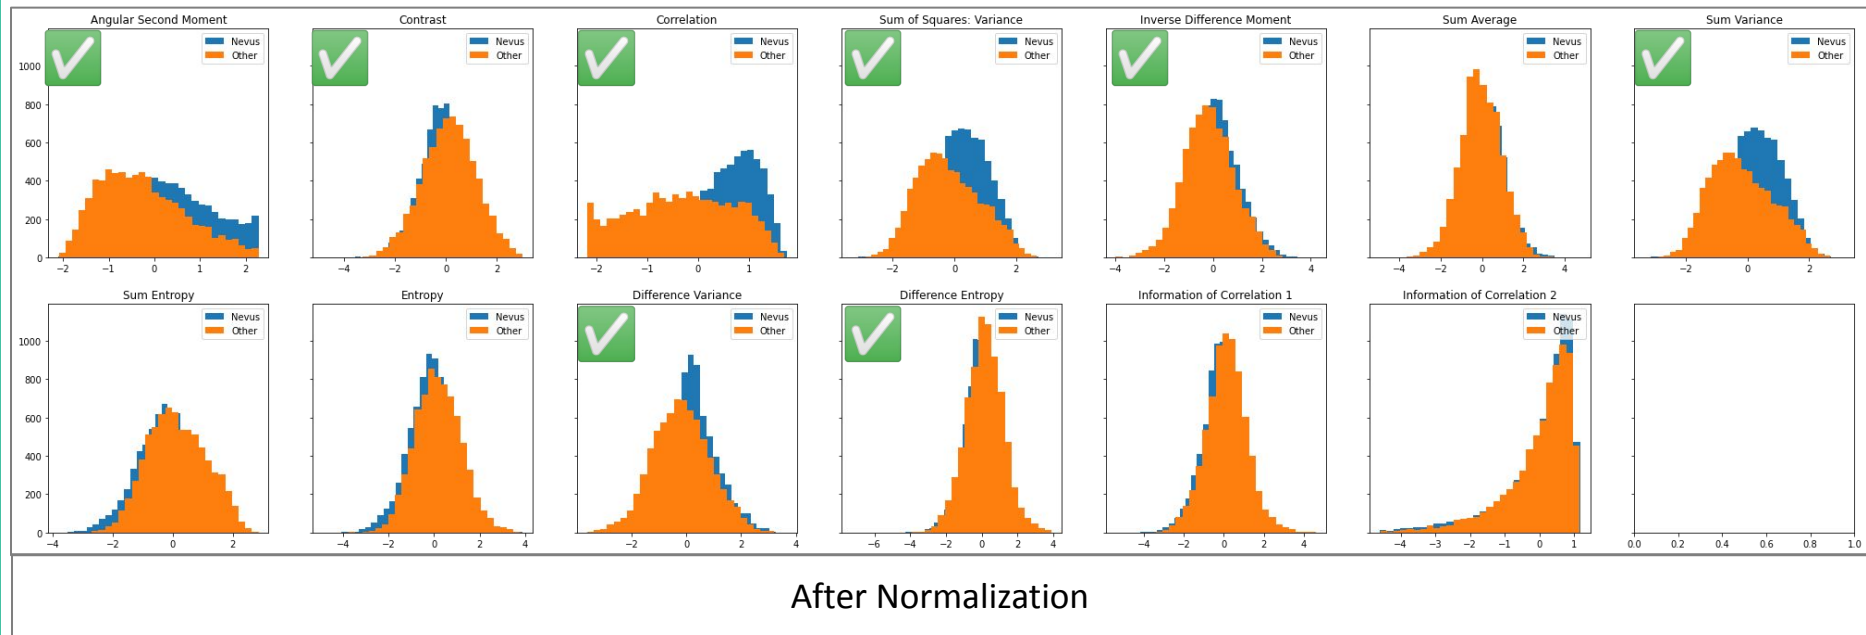
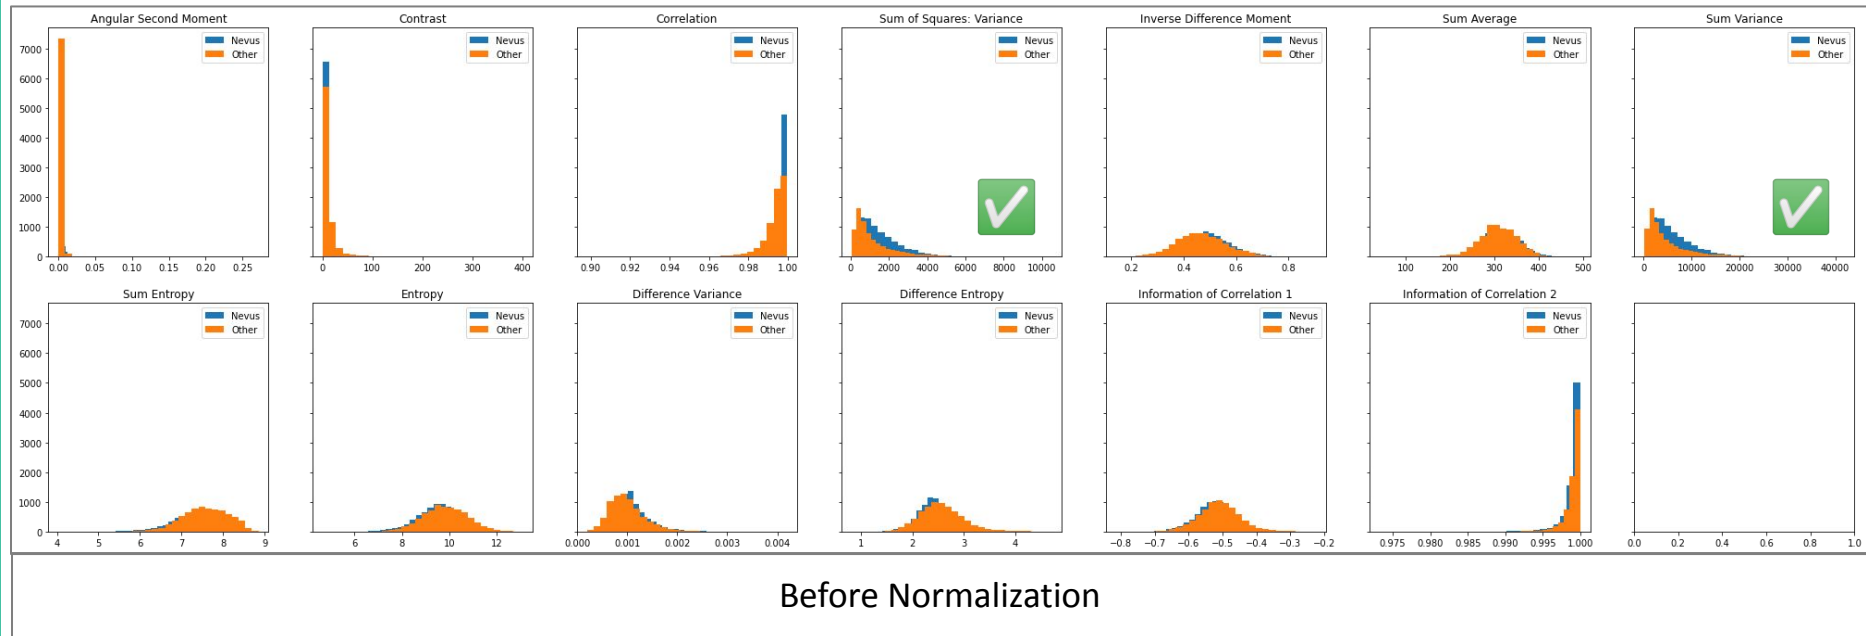


Before Normalization





After Normalization

# Haralick Features



# LBP Histograms




# Training and Evaluation

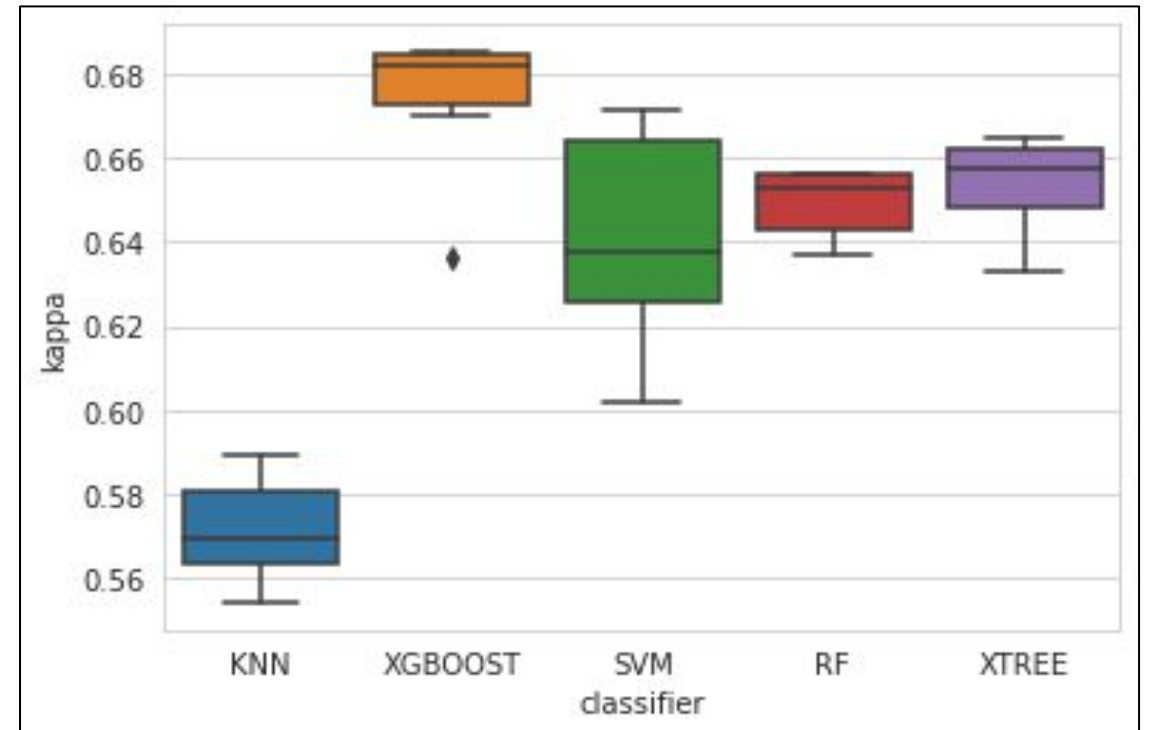
---

Feature selection & model training results

# Global Results Evaluation



Boxplot of the kappa metric for each feature vector across all models




Boxplot of the kappa metric for each classifier across all feature vectors




# Best Results Per Classifier

Classifier	Parameters	Features	Time	Accuracy	Kappa
K-NN	k=20	L	0035.718	0.794783	0.589447
XGboost	n_estimators = 2000, depth = 9, lr = 0.1	A	2214.809	0.842729	0.685325
SVM	C = 8	L	0318.782	0.835879	0.671699
RF	n_estims = 500, depth = 20, criterion = gini	A	0135.674	0.828240	0.656174
Extra Trees	n_estims = 2000, depth = 9, criterion = gini	J	0080.913	0.832718	0.665179
Stacking	XGboost + SVM + Extra Trees	L	10339.71	0.840358	0.680631
Averaging (sklearn)	XGboost + SVM + Extra Trees	L	2687.811	0.842992	0.685815
Majority Voting (Sklearn)	XGboost + SVM + Extra Trees	L	2564.247	0.843256	0.686305
<b>Averaging (manual)</b>	<b>XGboost + SVM</b>	<b>L</b>	<b>—</b>	<b>0.844573</b>	<b>0.689023</b>
Majority Voting (manual)	XGboost + SVM + Extra Trees	L	—	0.844046	0.687884


# Best Classifier Results




Examples of Correctly Classified Images



Examples of Wrongly Classified Images






Challenge 2: Melanoma VS BCC VS SCC

# Notes

---

- For this challenge, we follow the same proposed pipeline as in Challenge 1.
- The same preprocessing steps were applied.
- The same features were extracted.





Sampling


# Class imbalance with SMOTE and Undersampling

---

For Challenge 2, class imbalance is tackled by sampling: Oversampling the minority class (~1700 samples) with SMOTE, then randomly undersampling the other classes to that same amount.

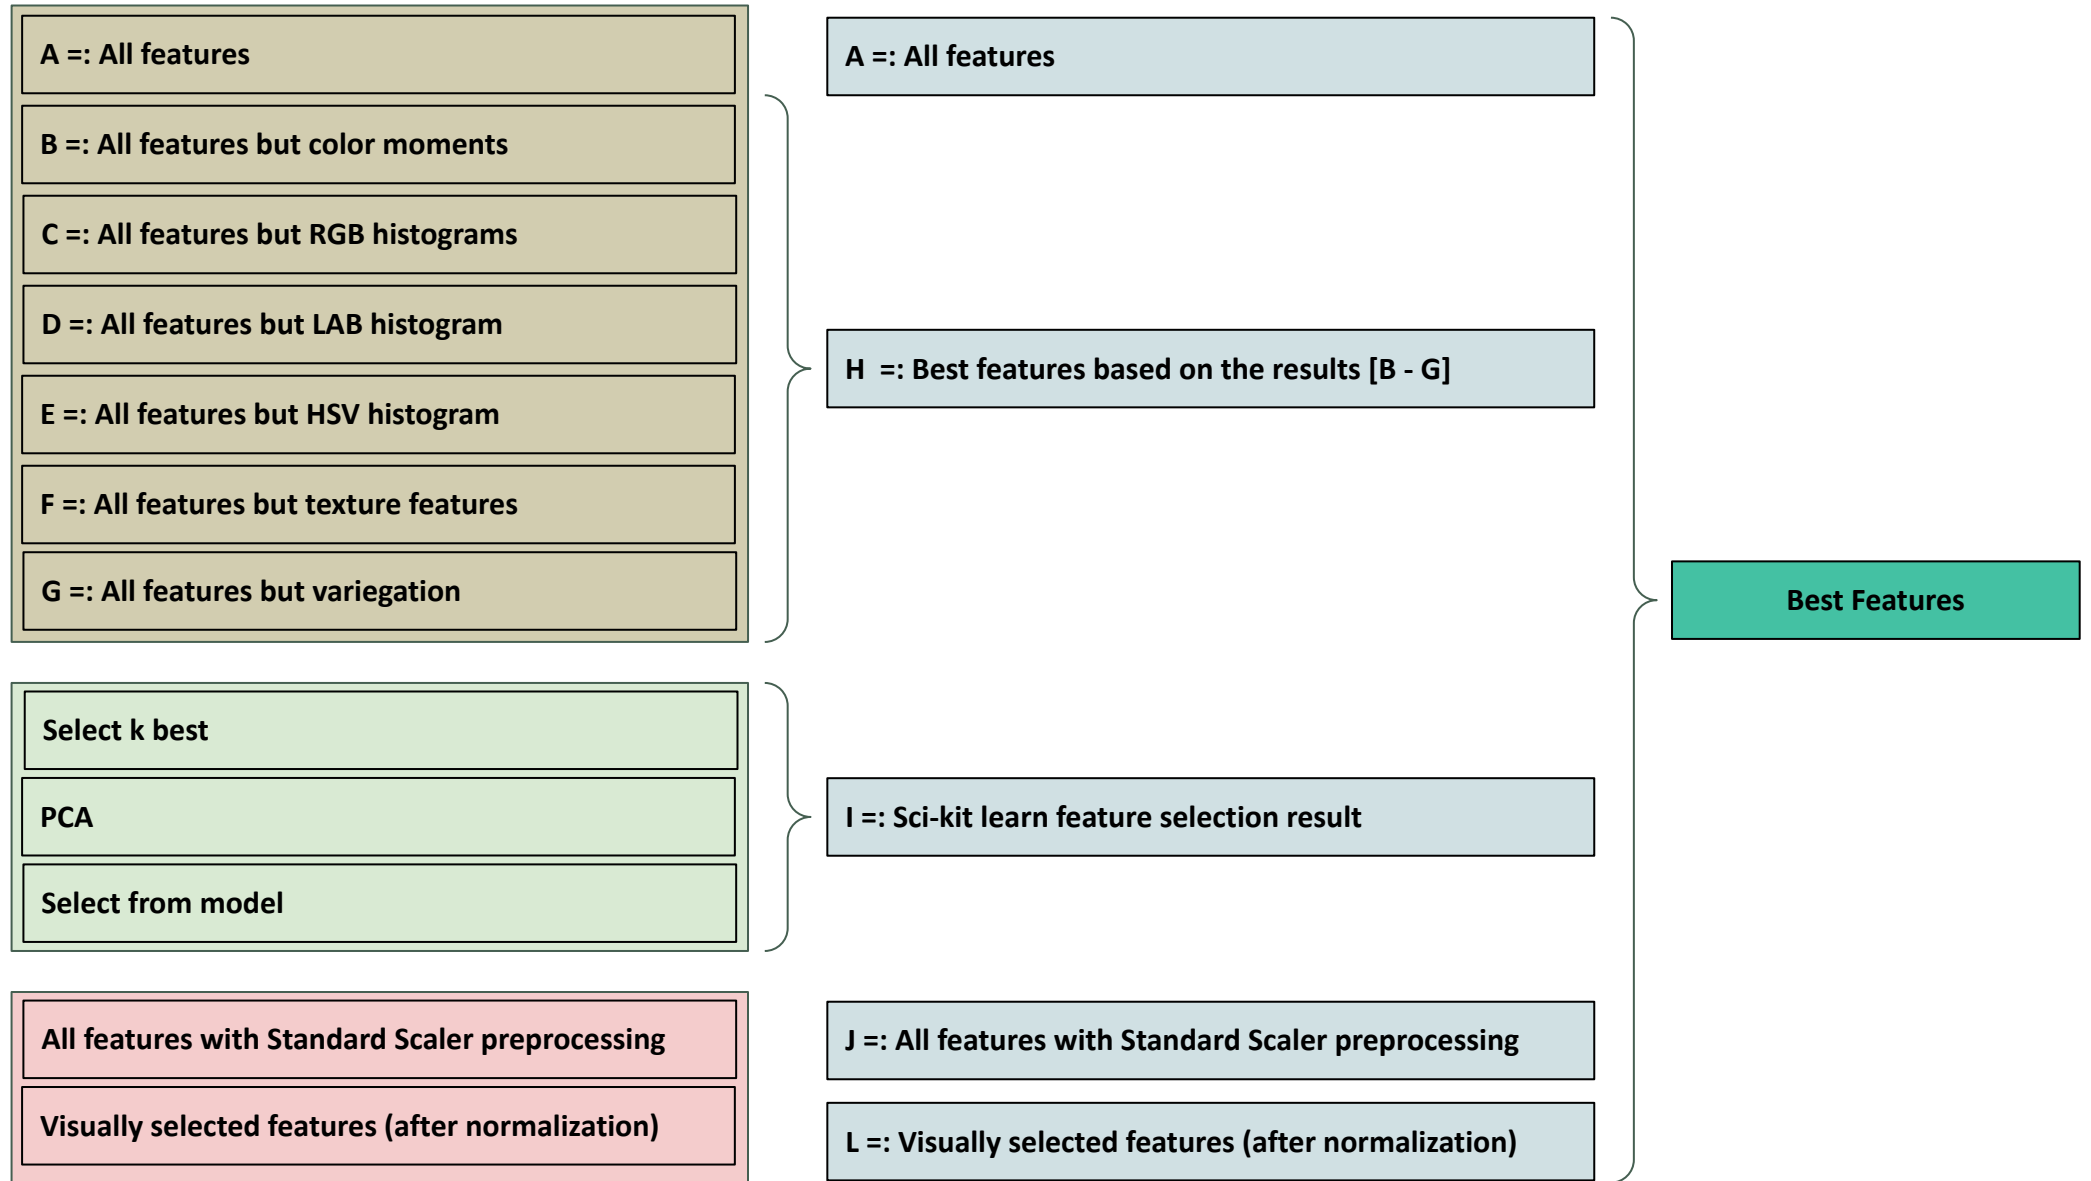
1. SMOTE : an oversampling technique used on minority class. It works by selecting samples from the minority class and finding the k-nearest neighbors of each of them. Then, synthetic examples are created at a randomly selected point between two of those neighbors.
2. Random Undersampling : Randomly selecting samples from the majority classes without replacement, while keeping the minority class unchanged.

## Synthetic Minority Oversampling Technique



# Feature Selection

**Feature Selection Pipeline (For each classifier)**





# Model Selection

# Challenge 2

The best model is selected based on the following evaluation metrics:

$$\text{Kappa} = \frac{\text{Pr}(a) - \text{Pr}(e)}{1 - \text{Pr}(e)}$$

$$\text{balanced-accuracy} = \frac{1}{3} \left( \frac{TP}{TP + FN} \right)$$


		Predicted Labels		
		Mel	Bcc	Scc
True Labels	Scc	TN	FP	TN
	Bcc	FN	TP	FN
	Mel	TN	FP	TN

# Challenge 2 (Cont'd)

---

The following models were compared:

- Nearest Neighbor (Baseline)
- K-Nearest Neighbors
- Support Vector Machines
- Gradient Boosting (XGBoost)
- Extra Trees
- Random Forest
- Stacking ensemble
- Majority voting ensemble (soft)
- Ensembling the pre-trained models with averaging
- Hierarchical architecture with previous classifiers





Experimental Results

# Feature Visualization

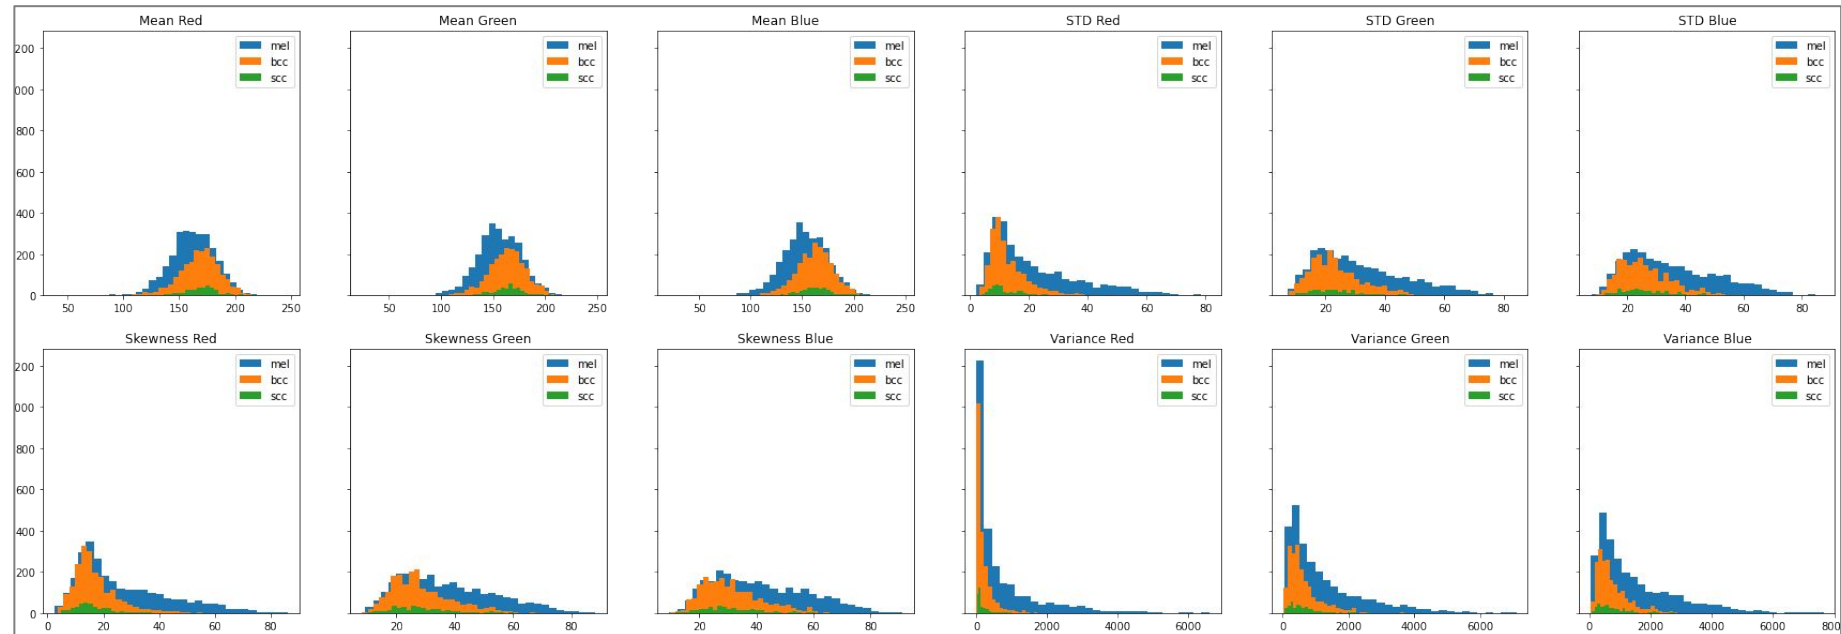
---

With Matplotlib

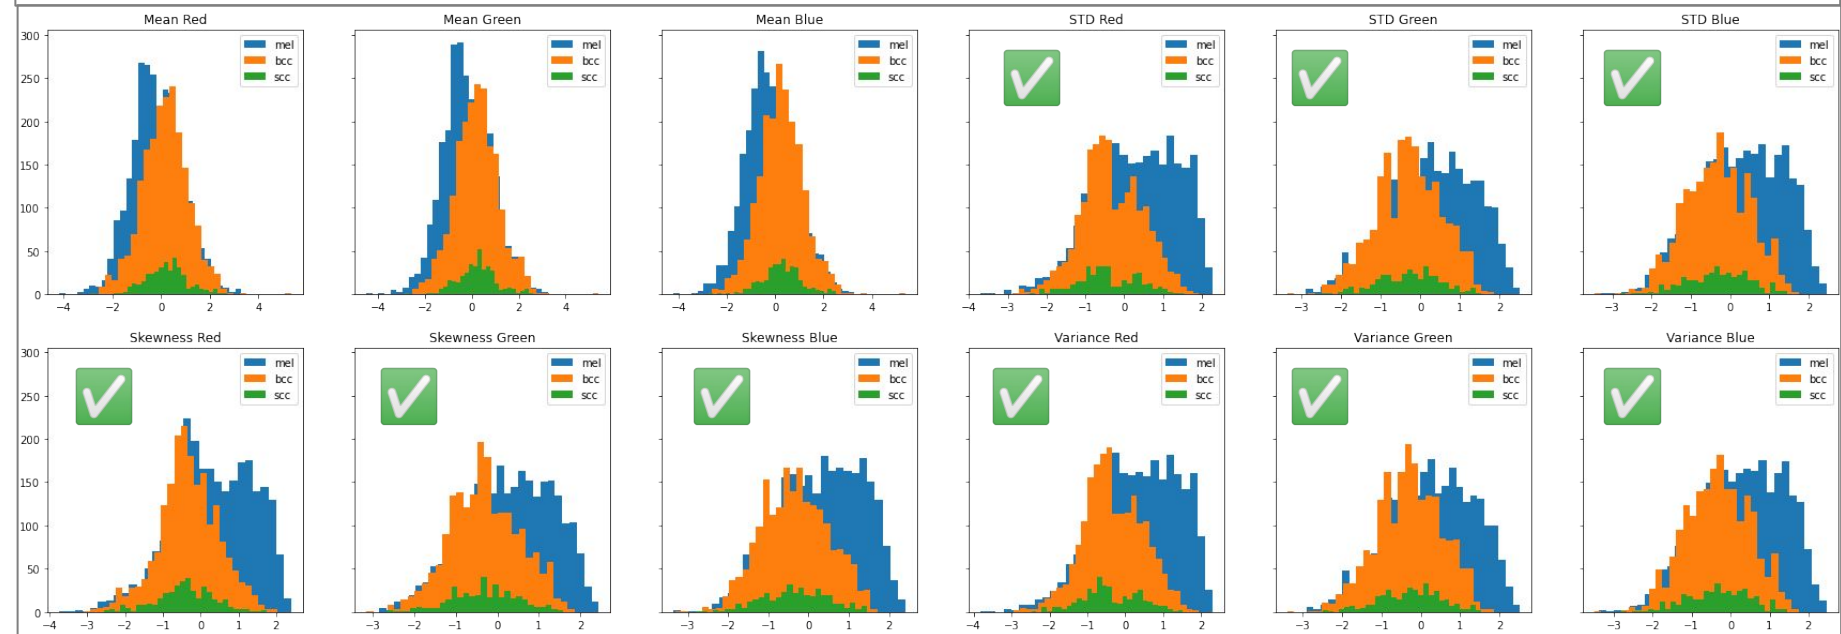
# Variegation



# Color Moments





Before Normalization





After Normalization

# Mean RGB Histograms For Each Class







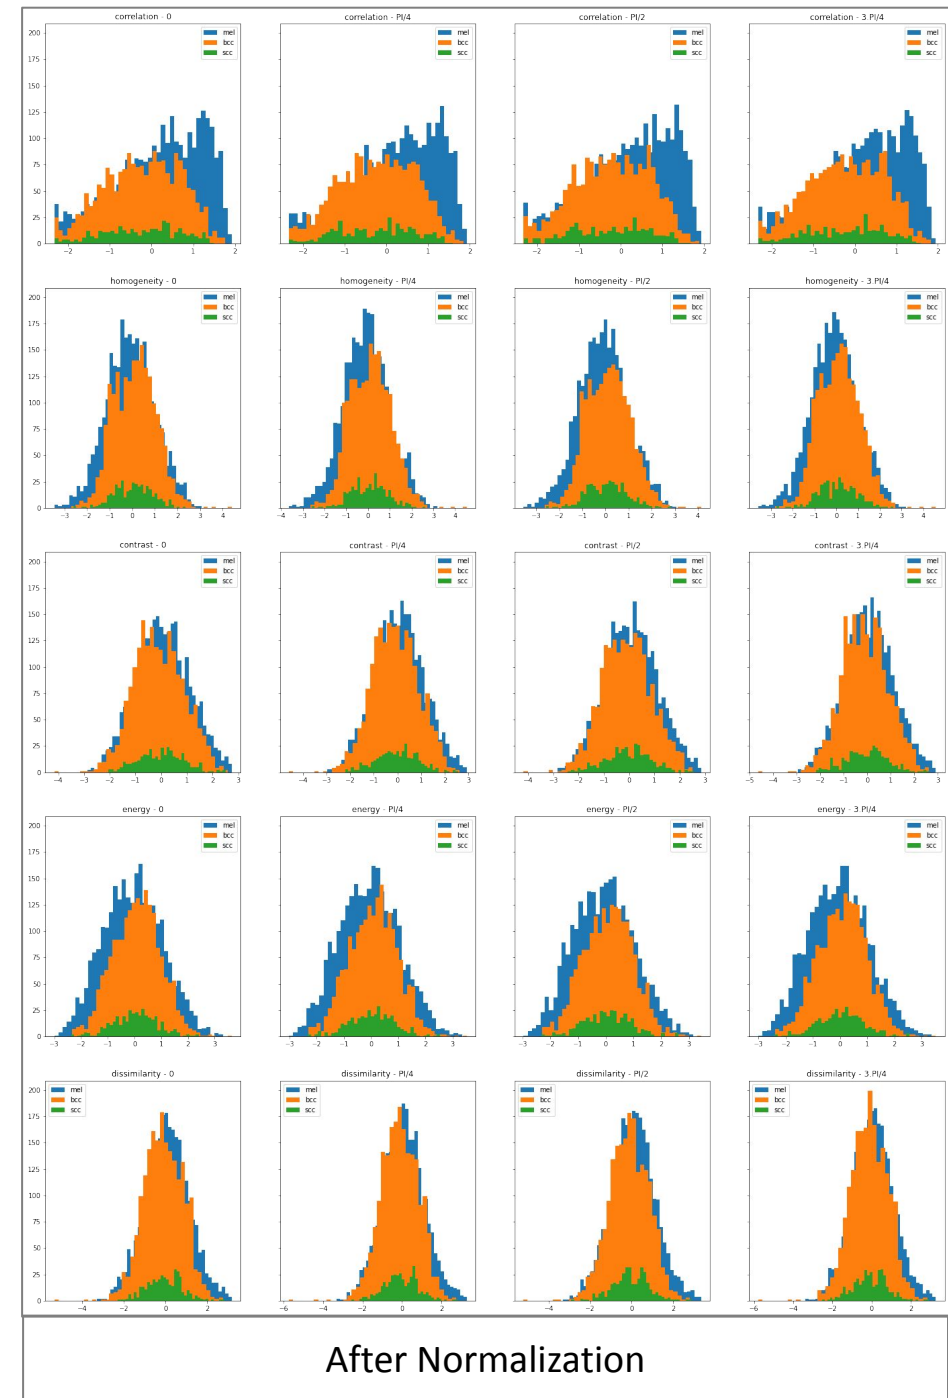
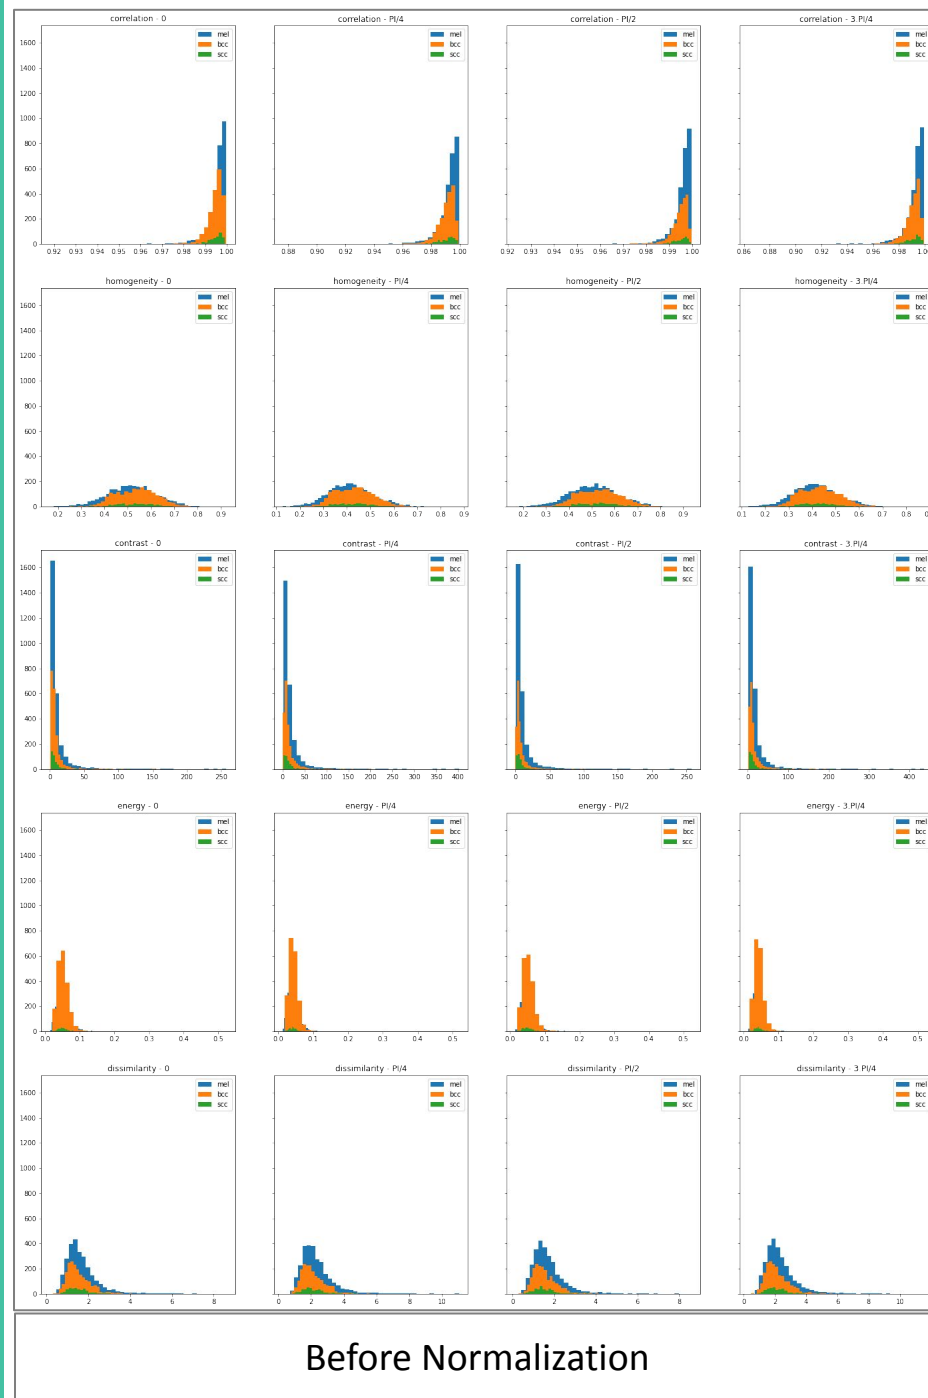
# Mean HSV Histograms For Each Class



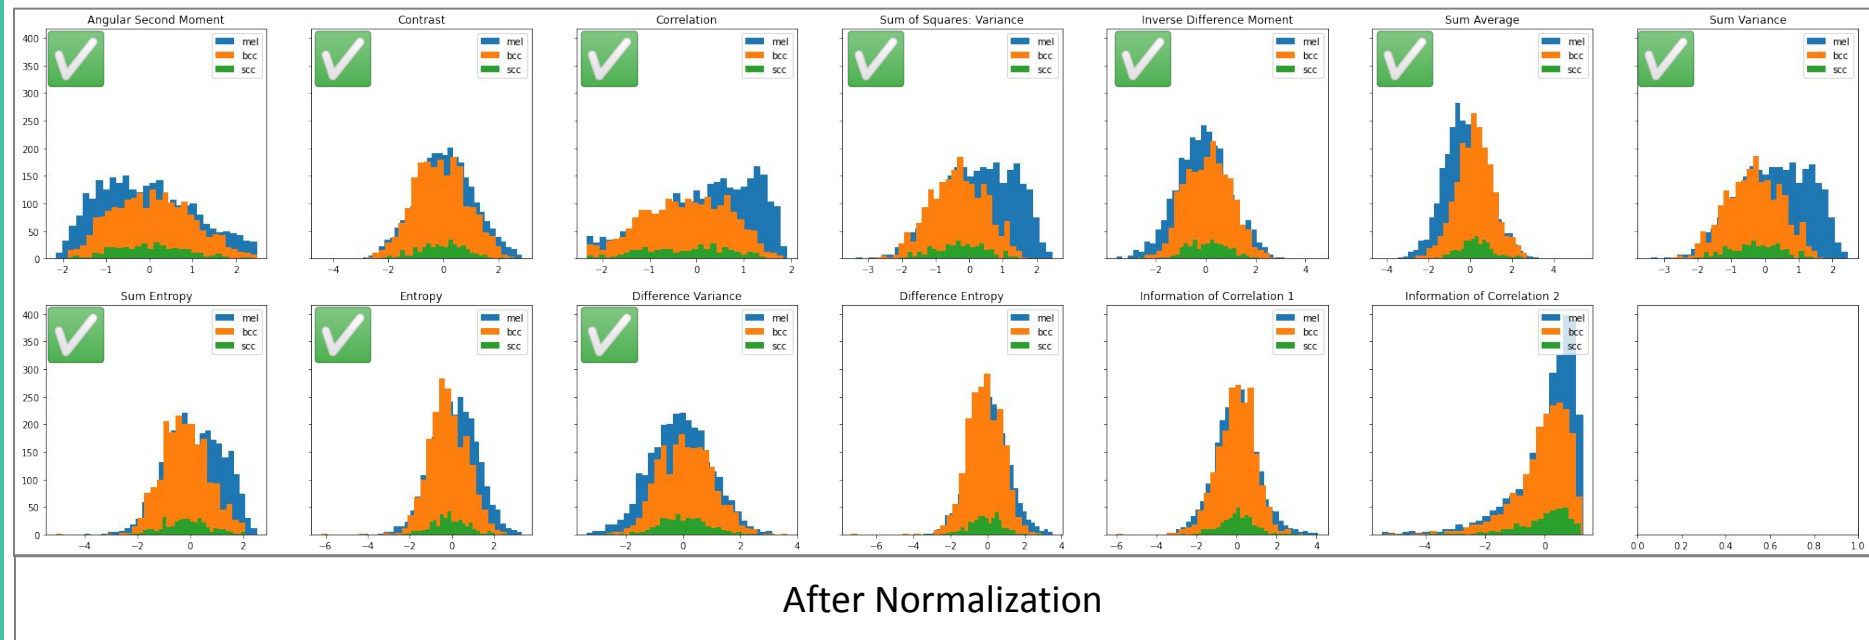
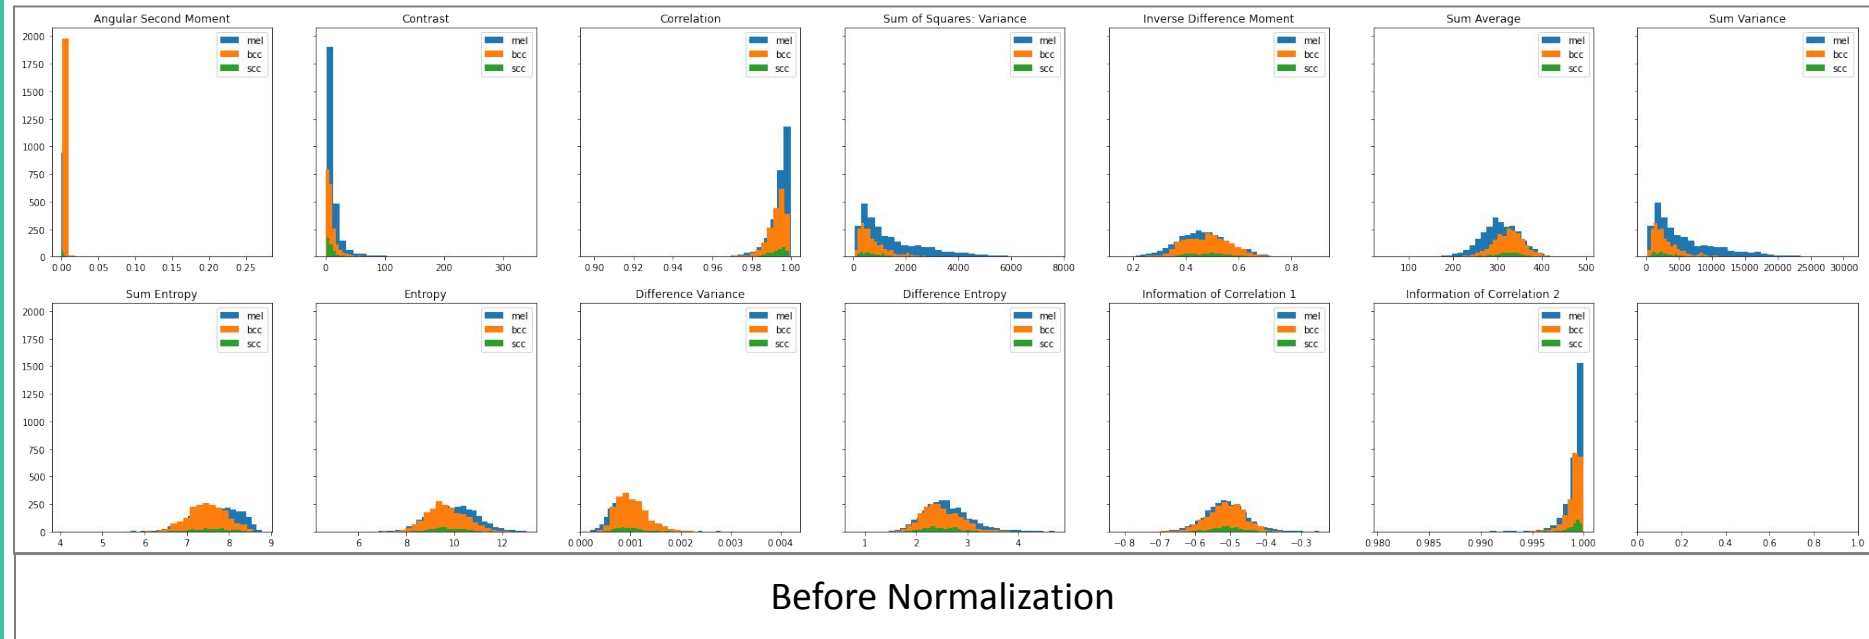
# Mean LAB Histograms For Each Class





# GLCM Features



# Haralick Features



# LBP Histograms




# Training and Evaluation


---

Feature selection & model training results

# Global Results Evaluation



Boxplot of the kappa metric for each feature vector across all models




Boxplot of the kappa metric for each classifier across all feature vectors




# Best Results Per Classifier

Classifier	Parameters	Features	Kappa	Bma	Time (s)
Baseline NN	k =1	A	0.621365	0.747204	17.2943
K-NN	k=5	I	0.593095	0.742710	14.6484
XGboost	n_estimators = 500, depth = 7, lr = 0.1	E	0.693371	0.721120	419.1023
SVM	C = 11	I	0.735959	0.747785	635.5958
RF	n_estimators = 500, depth = 13, criterion = entropy	A	0.634628	0.698730	99.1131
Extra Trees	n_estimators = 1000, depth = 13, criterion = entropy	G	0.639188	0.716365	50.6742
Stacking (sklearn)	SVM + XGBoost + XTrees + RF + KNN	A	0.734990	0.741193	2300.3030
<b>Majority Voting (Sklearn)</b>	<b>SVM + XGBoost</b>	<b>A</b>	<b>0.752323</b>	<b>0.741561</b>	<b>497.1405</b>
Averaging (manual)	SVM + XGBoost + XTrees	A	0.731742	0.742417	-
Hierarchical (manual)	XGBoost	A	0.652815	0.674347	-



# Best Classifier Results




Examples of Correctly Classified Images



Examples of Wrongly Classified Images

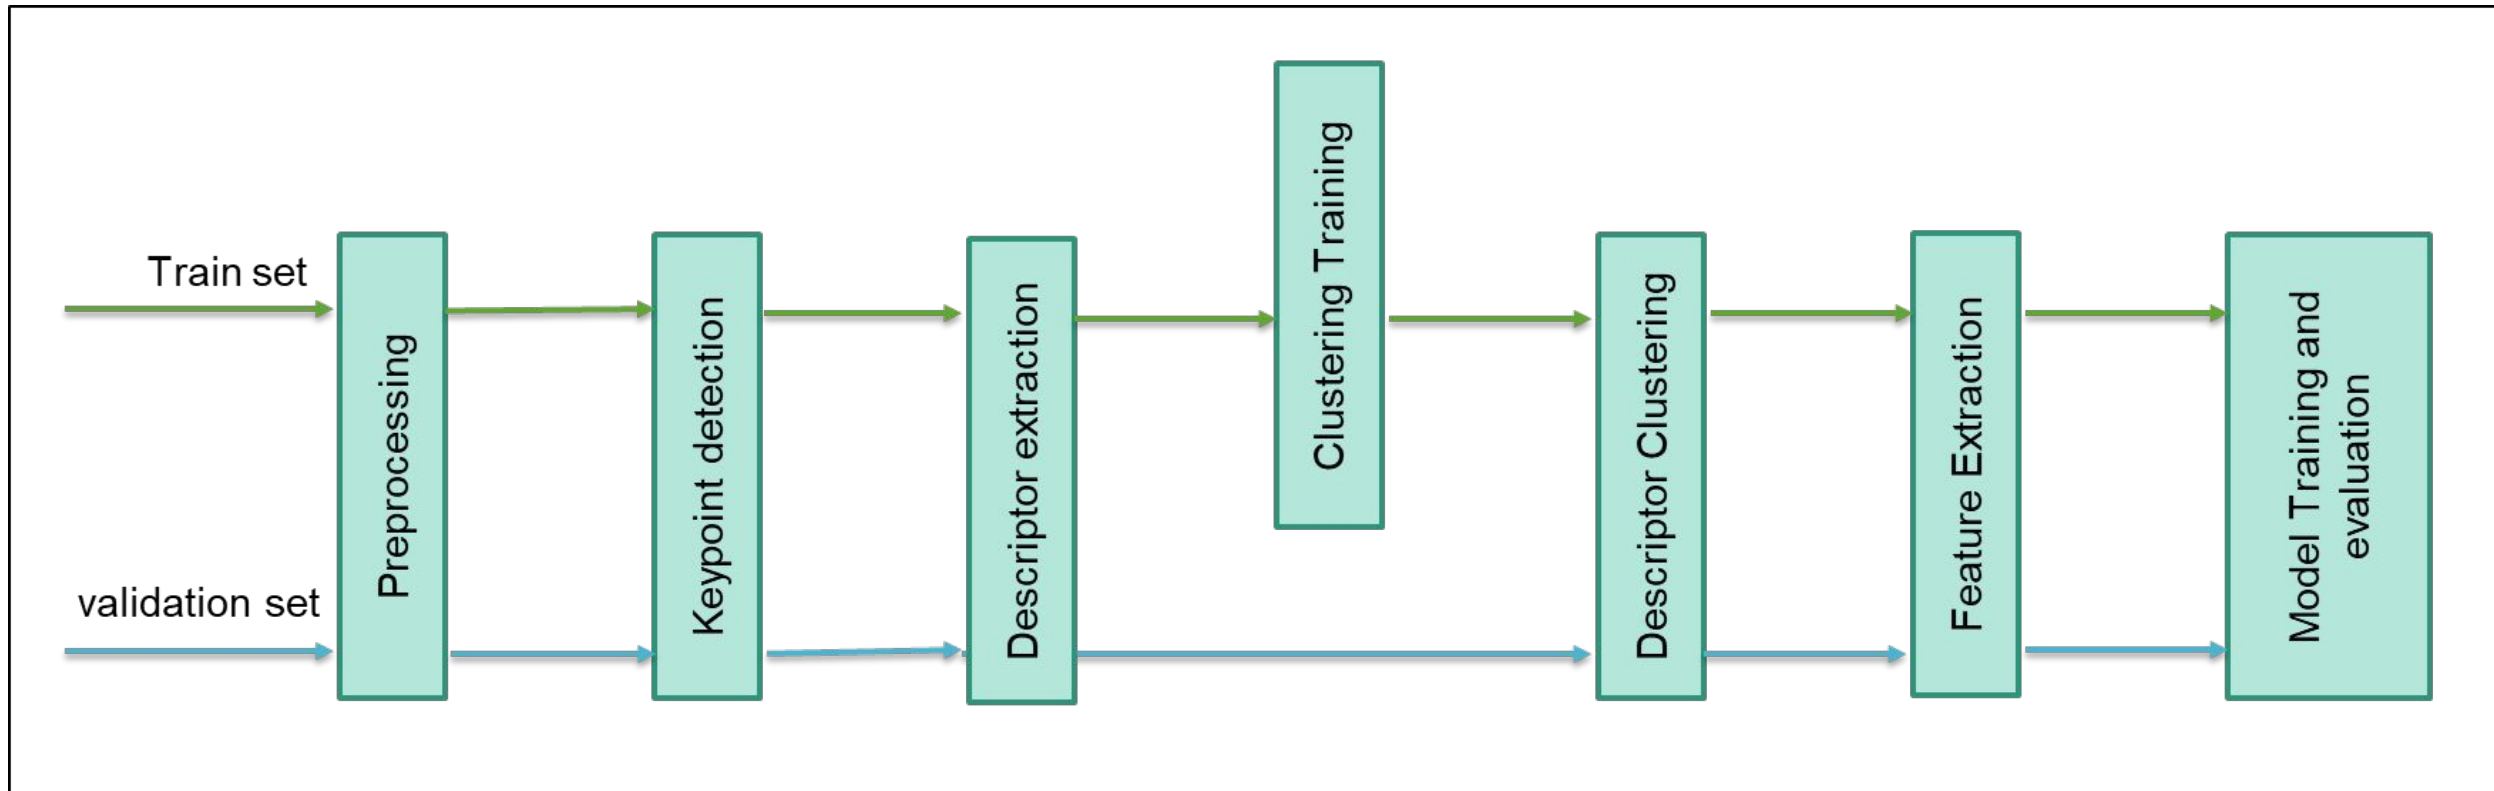


Confusion matrix on the validation set



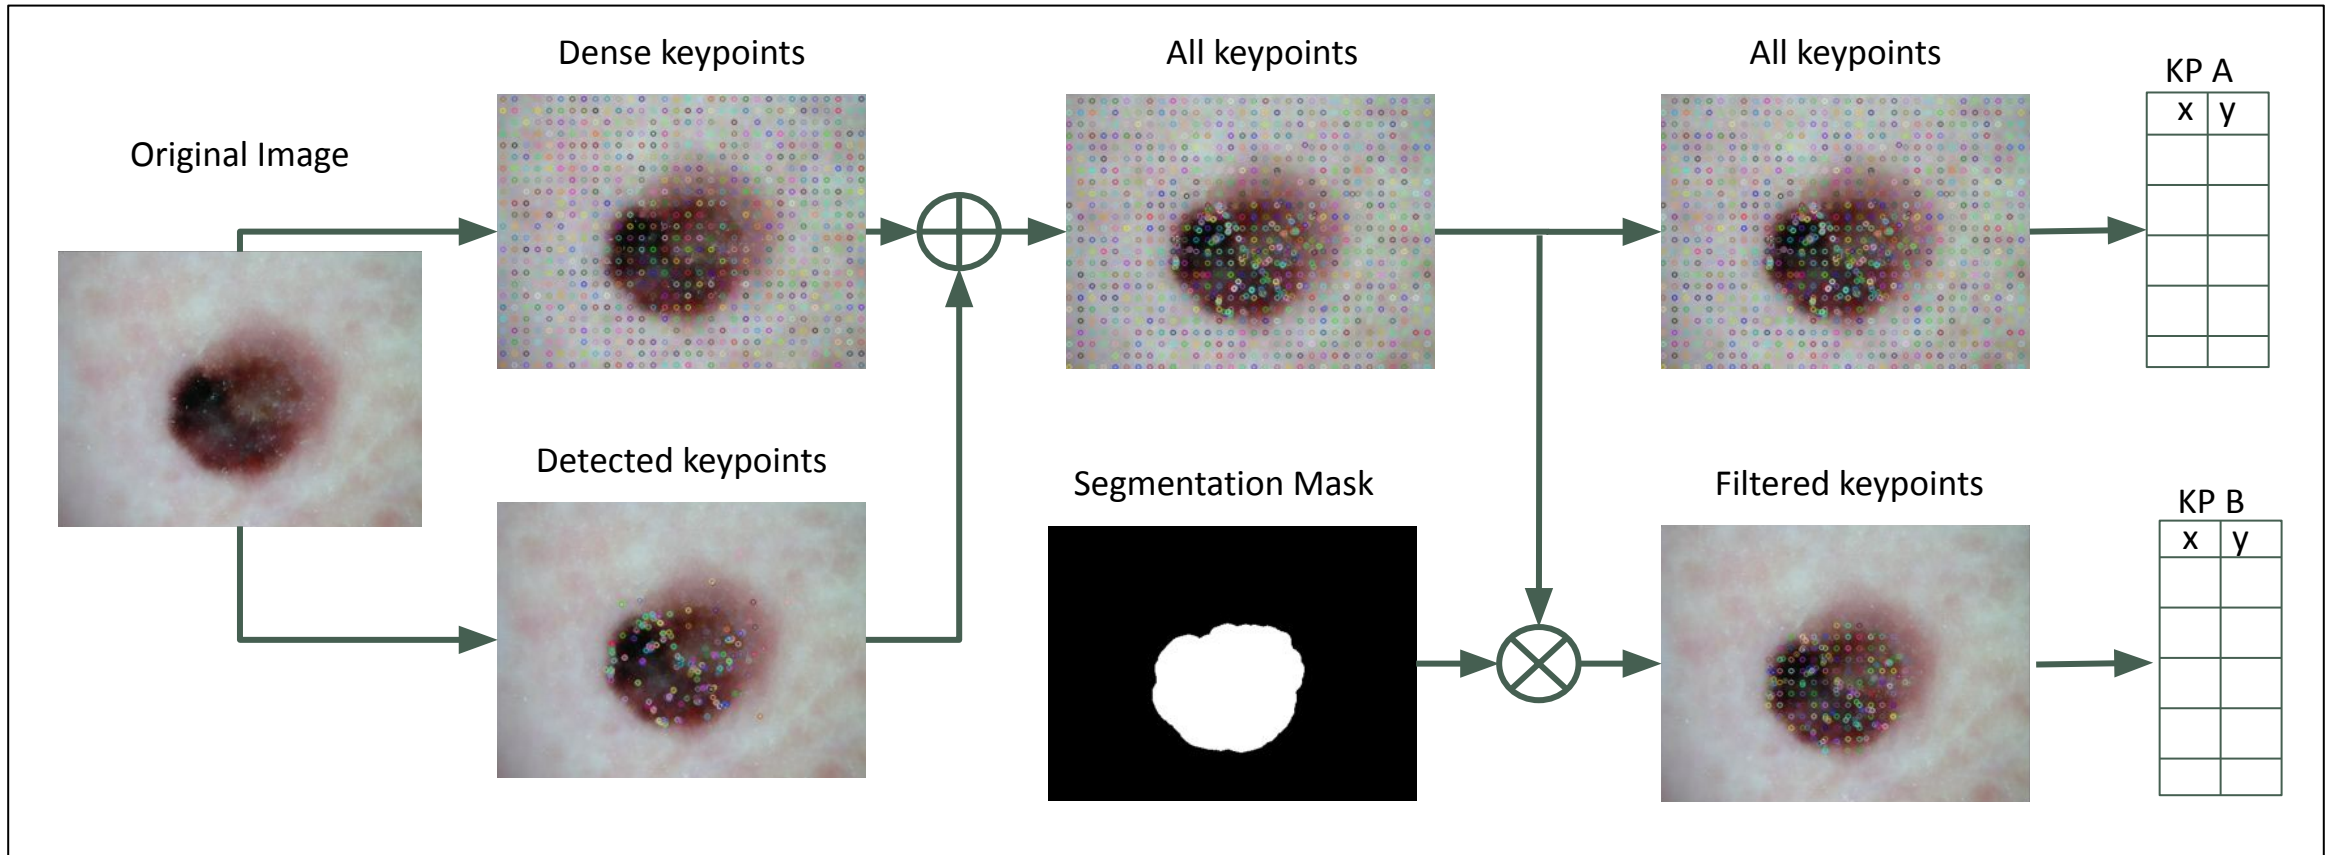
Bag Of Visual Words

# Pipeline




Bag-of-visual-words based classification pipeline

# Keypoint detection



Keypoint Detection Pipeline

# Segmentation

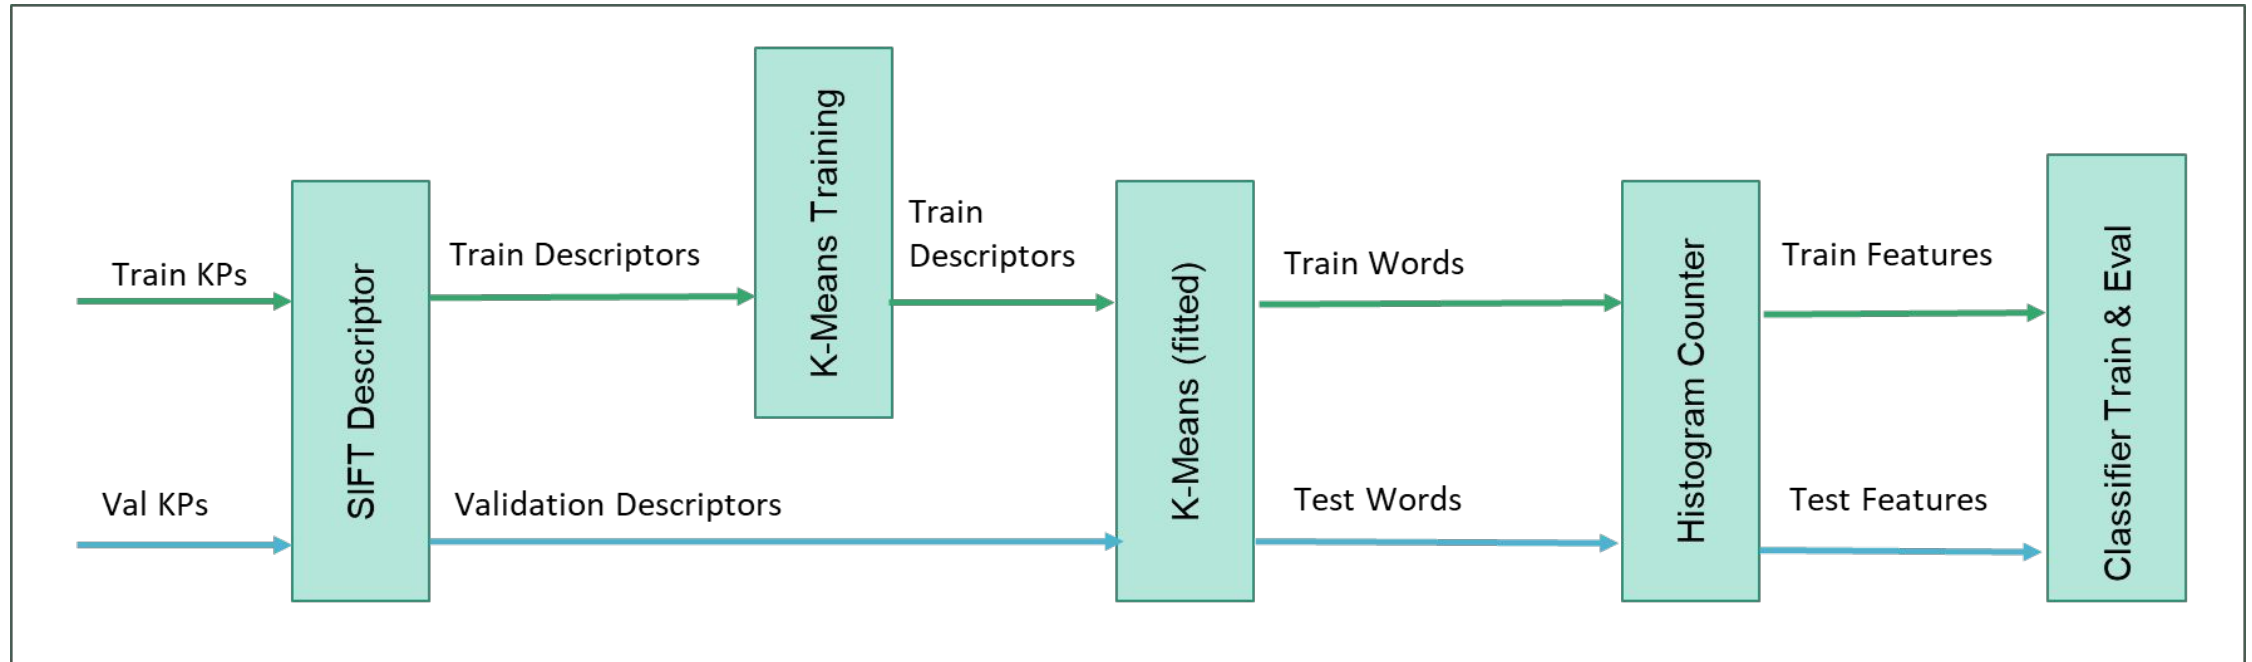


Segmentation pipeline

\*ASF : Alternating Sequential Filtering to reduce image complexity and make it more homogeneous

\*Mask : removing some artefacts (colorful disks) on the borders when using the blue channel


# Feature Extraction



Feature Extraction Pipeline



# Extracted Features



Histogram of the vocabulary occurrences in the train set in each class using SIFT descriptor.

# Model Training and Evaluation

---

- GridSearchCV with 5 folds was used for parameter estimation.
- SMOTE and undersampling were used for balancing the dataset.
- The following models were evaluated:
  - K-NN
  - SVM
  - XGBoost
  - Random Forest
  - Extra Trees
  - Ensembles of the previously trained models
- The models were compared based on balanced multi-class accuracy and kappa metrics.
- Models were trained using a combination of SIFT + BRISK + ORB features.

# Experimental Results

---

With Segmentation mask


Model	Train Data	Best Kappa	BMA
NN	All samples	0.486241	0.595114
KNN	All samples	0.486241	0.595114
<b>XGBST</b>	<b>1694 samples/class</b>	<b>0.489844</b>	<b>0.560988</b>
SVM	All samples + balanced weights	0.457248	0.588034
RF	All samples + balanced weights	0.478247	0.516117
XTrees	1694 samples/class	0.472977	0.562755


Without segmentation mask

Model	Train Data	Best Kappa	BMA
NN	All samples	0.530192	0.646191
KNN	All samples	0.530192	0.646191
<b>XGBST</b>	<b>1694 samples/class</b>	<b>0.545448</b>	<b>0.601006</b>
SVM	1694 samples/class	0.537877	0.656429
RF	All samples + balanced weights	0.515091	0.543635
XTrees	All samples	0.543117	0.601907

# Experimental Results: SIFT+ORB+BRISK

Model	Train Data	Best Kappa	BMA
NN	All samples	0.531296	0.661191
KNN	All samples	0.531296	0.661191
XGBST	All samples	0.579542	0.574249
SVM	1694 samples/class	0.578807	0.637399
RF	All samples	0.524012	0.535706
XTrees	1694 samples/class	0.536174	0.589971
<b>Ensemble</b>	<b>All data + all models</b>	<b>0.599591</b>	<b>0.601939</b>





Conclusion

# Conclusion

- For skin lesion classification, Color and texture features play a major role in distinguishing classes.
- Sampling helps tackling imbalanced data classes.
- Normalized features to a gaussian-like distribution minimizes skewness and improves results.
- Ensembling classifiers improves the performance.
- SVM and XGB are indeed robust and they provide the best results.

# Future prospects

- Apply segmentation to enhance the ROI and focus on relevant features only.
- Improve the Hierarchical approach.
- Optimize the BoW approach by tuning the vocabulary size and adding more descriptors (e.g. color)

Thank you

---